

Last Lecture

Eugene Wu

Administrivia

AA4: git pull for bug fix

We will present project I highlights Thursday

Midterm review Friday 10AM location TBA

Scribenotes open until day of midterm

Midterm next Tuesday 8:30

Same location as midterm I

1 page 8x11 cheat sheet, both sides

Phone-a-friend

DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc
Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, etc

Key-value: Mongo, BerkeleyDB, Cassandra, etc

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc
Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, etc

Key-value: Mongo, BerkeleyDB, Cassandra, etc

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

Modern Database Systems

Hardware changes affect

Compression is good → Column stores

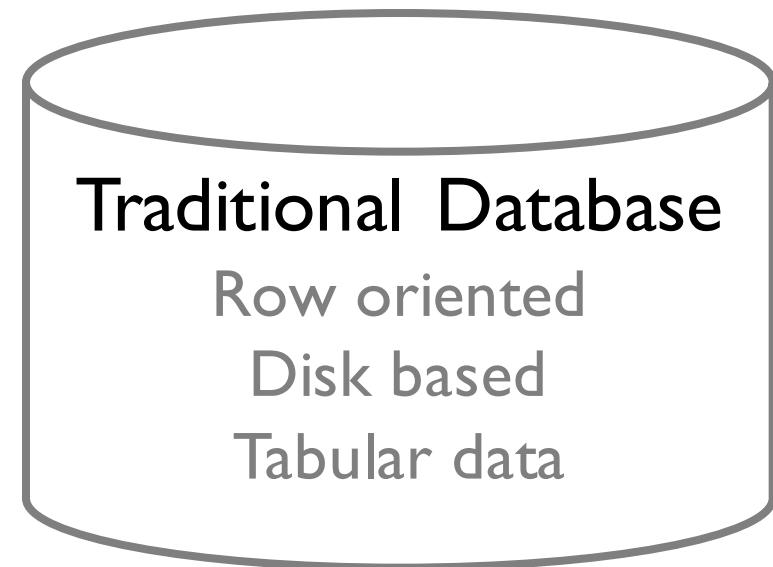
Large scale aggregation queries

“Data Warehouses”

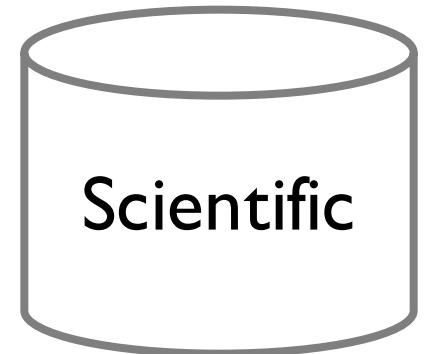
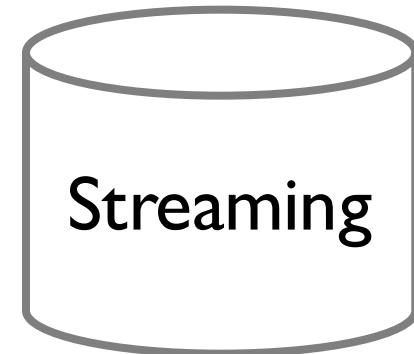
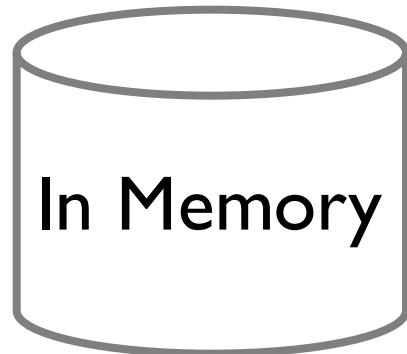
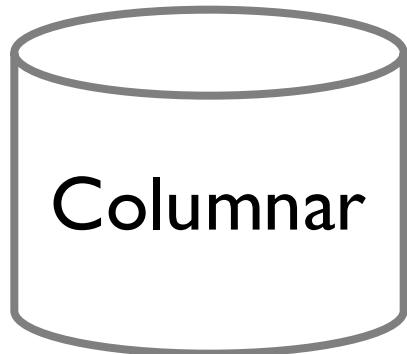
Memory is cheap → In-memory stores

Transactional systems

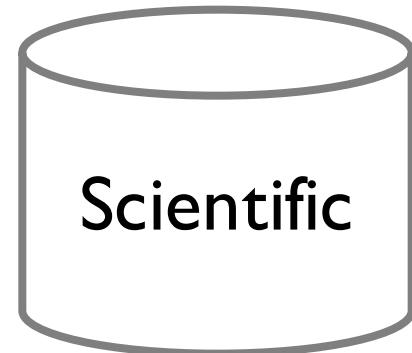
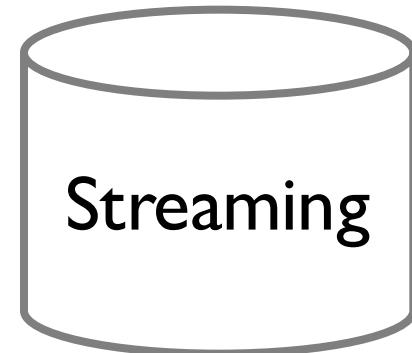
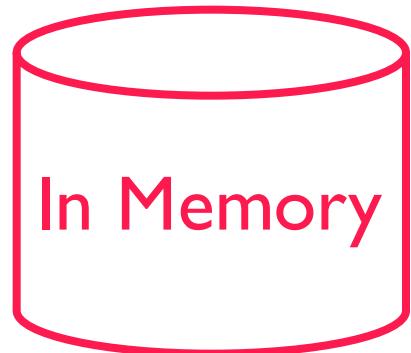
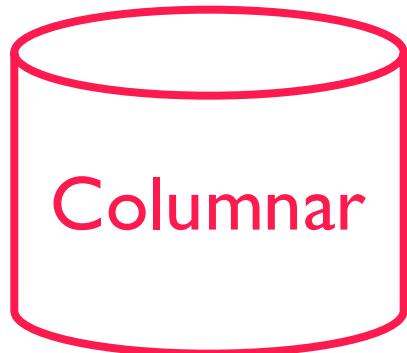
One Size Fits All



One Size Does Not Fit All



One Size Does Not Fits All



Data Warehouses

Store all historical data for future analysis

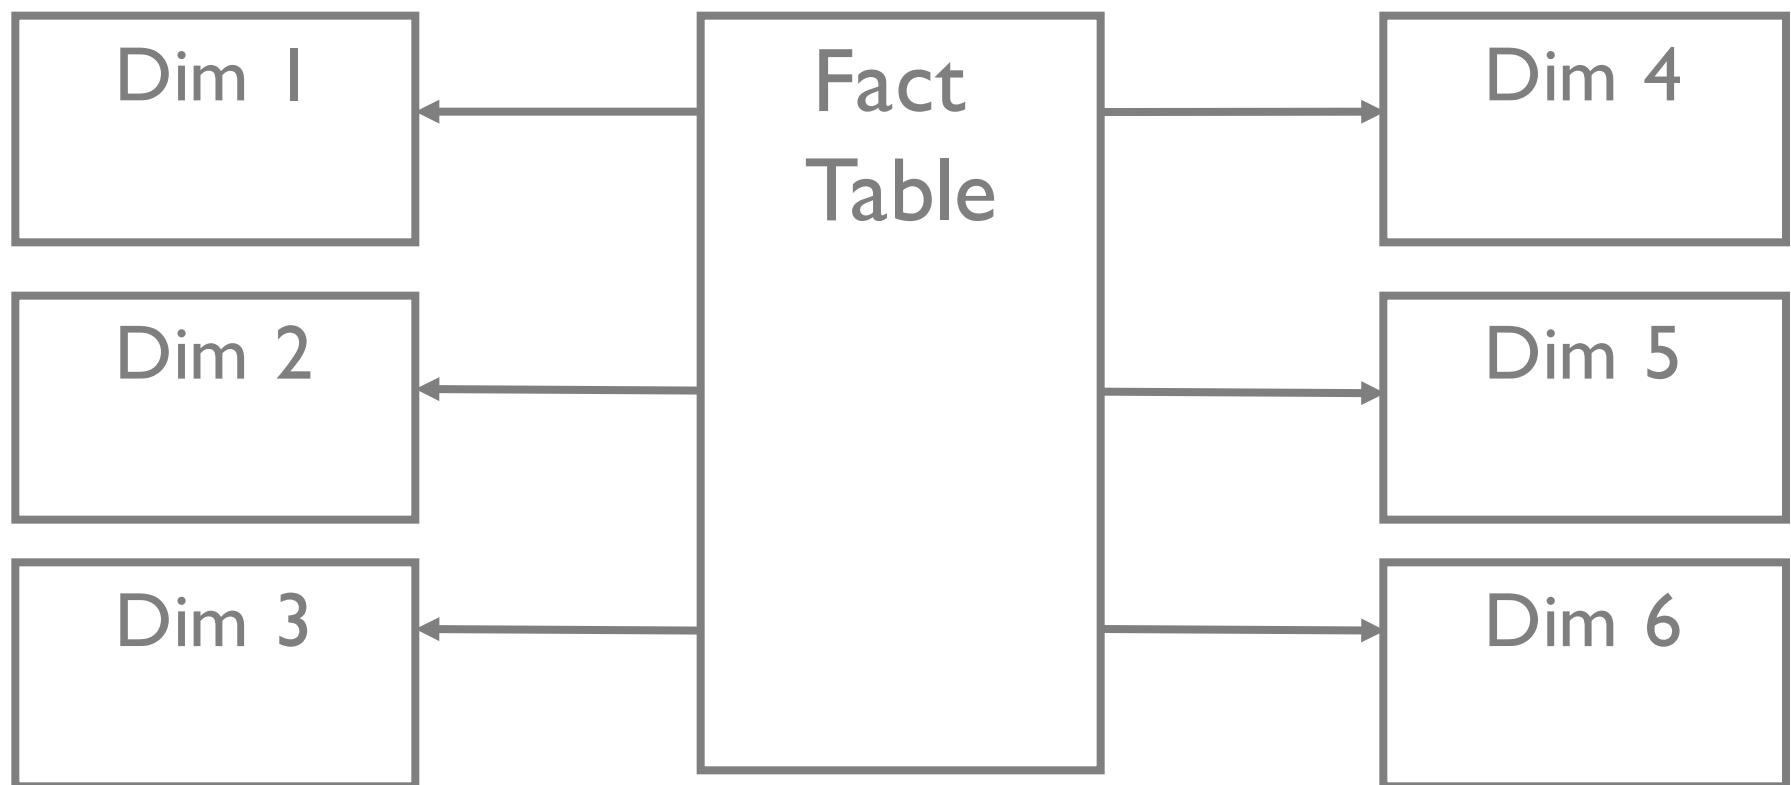
Sales by month over past 20 years

Clicks by youth in texas

Cost by product component

Defacto standard at any company

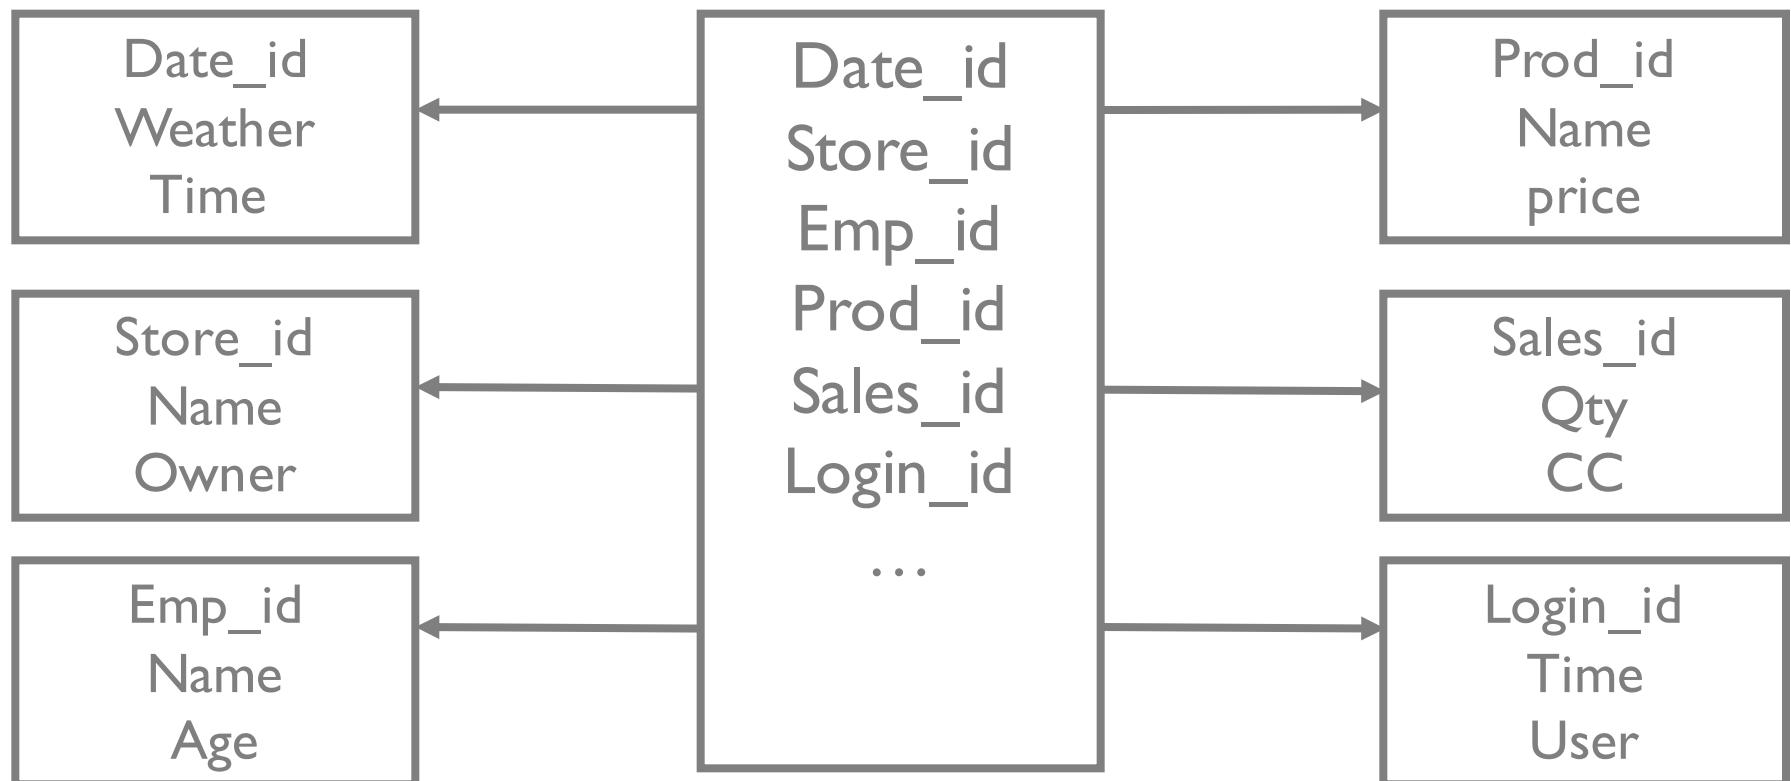
Star Schema



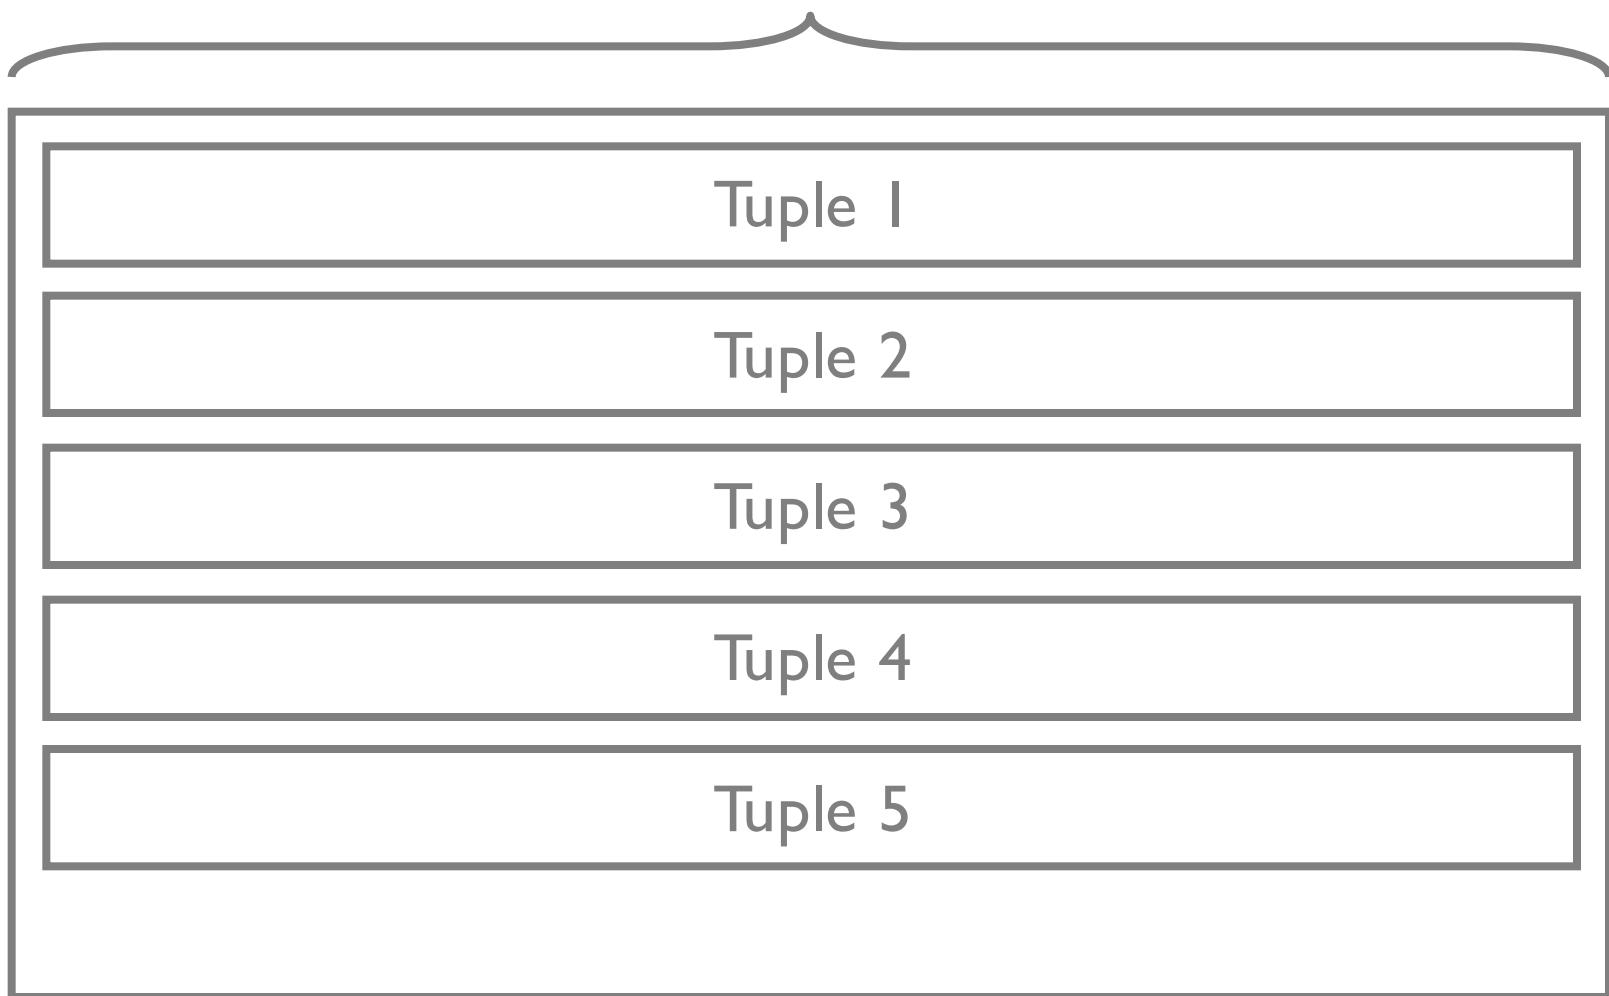
Star Schema

Fact table is “fat”

Queries access ~6 attrs



100 attributes



A1

A2

A3

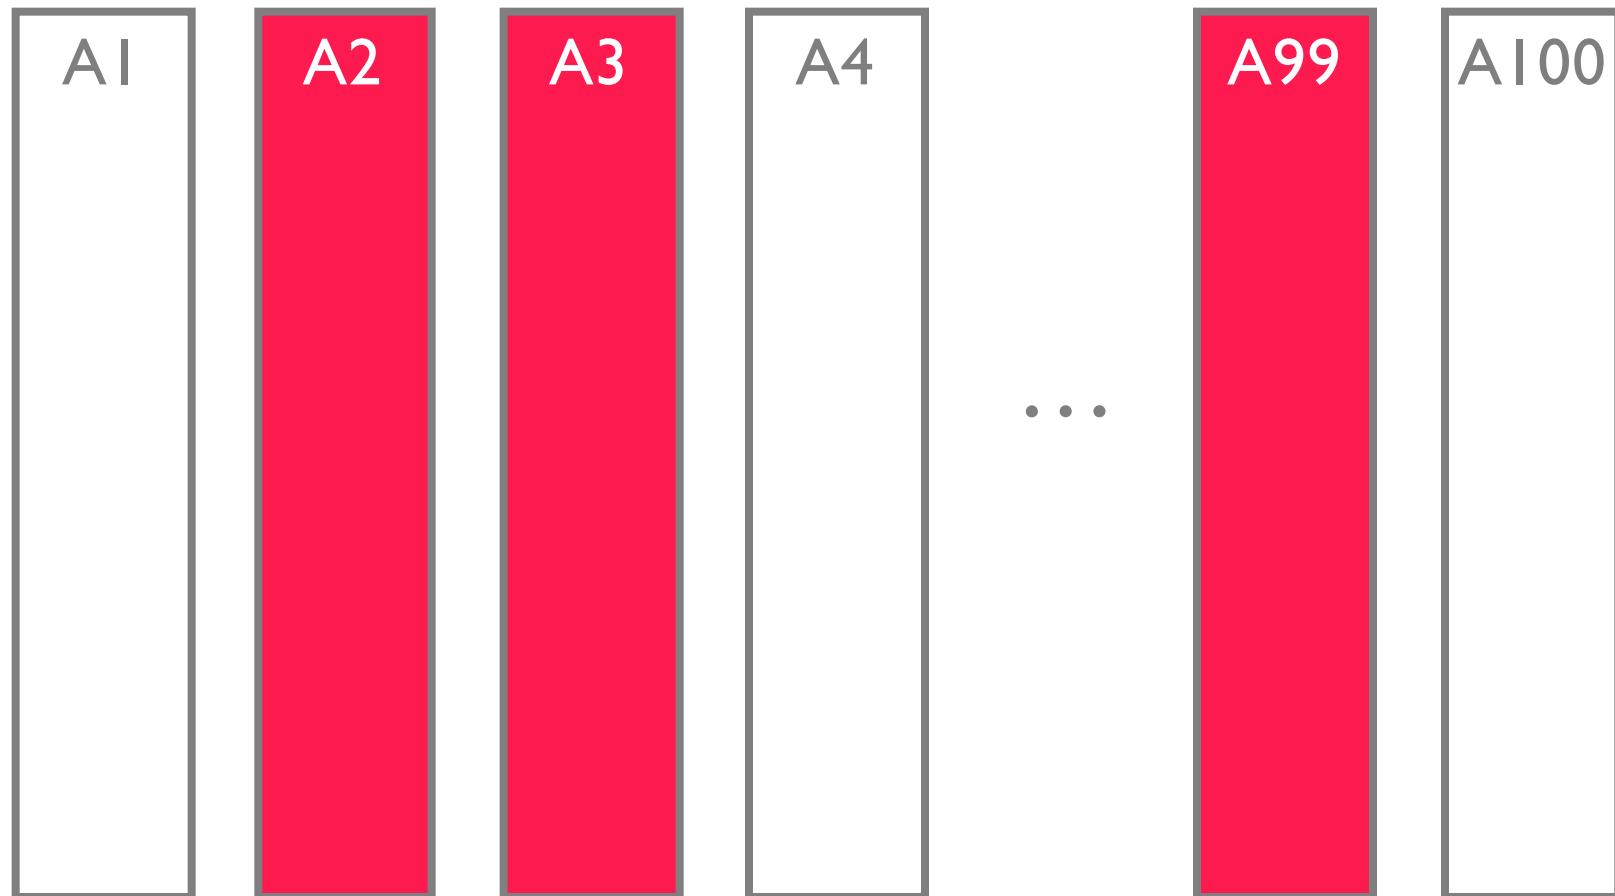
A4

...

A99

A100

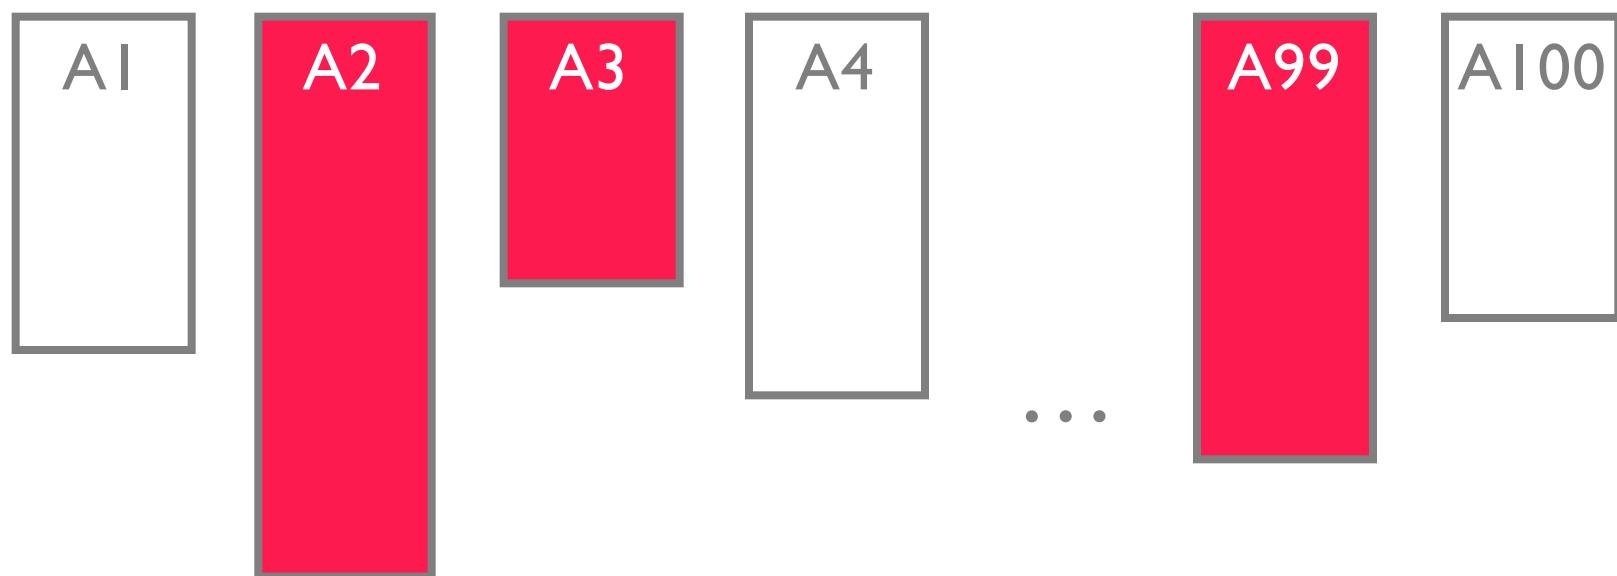
16x less data read. Unfair advantage.



16x less data read. Unfair advantage.

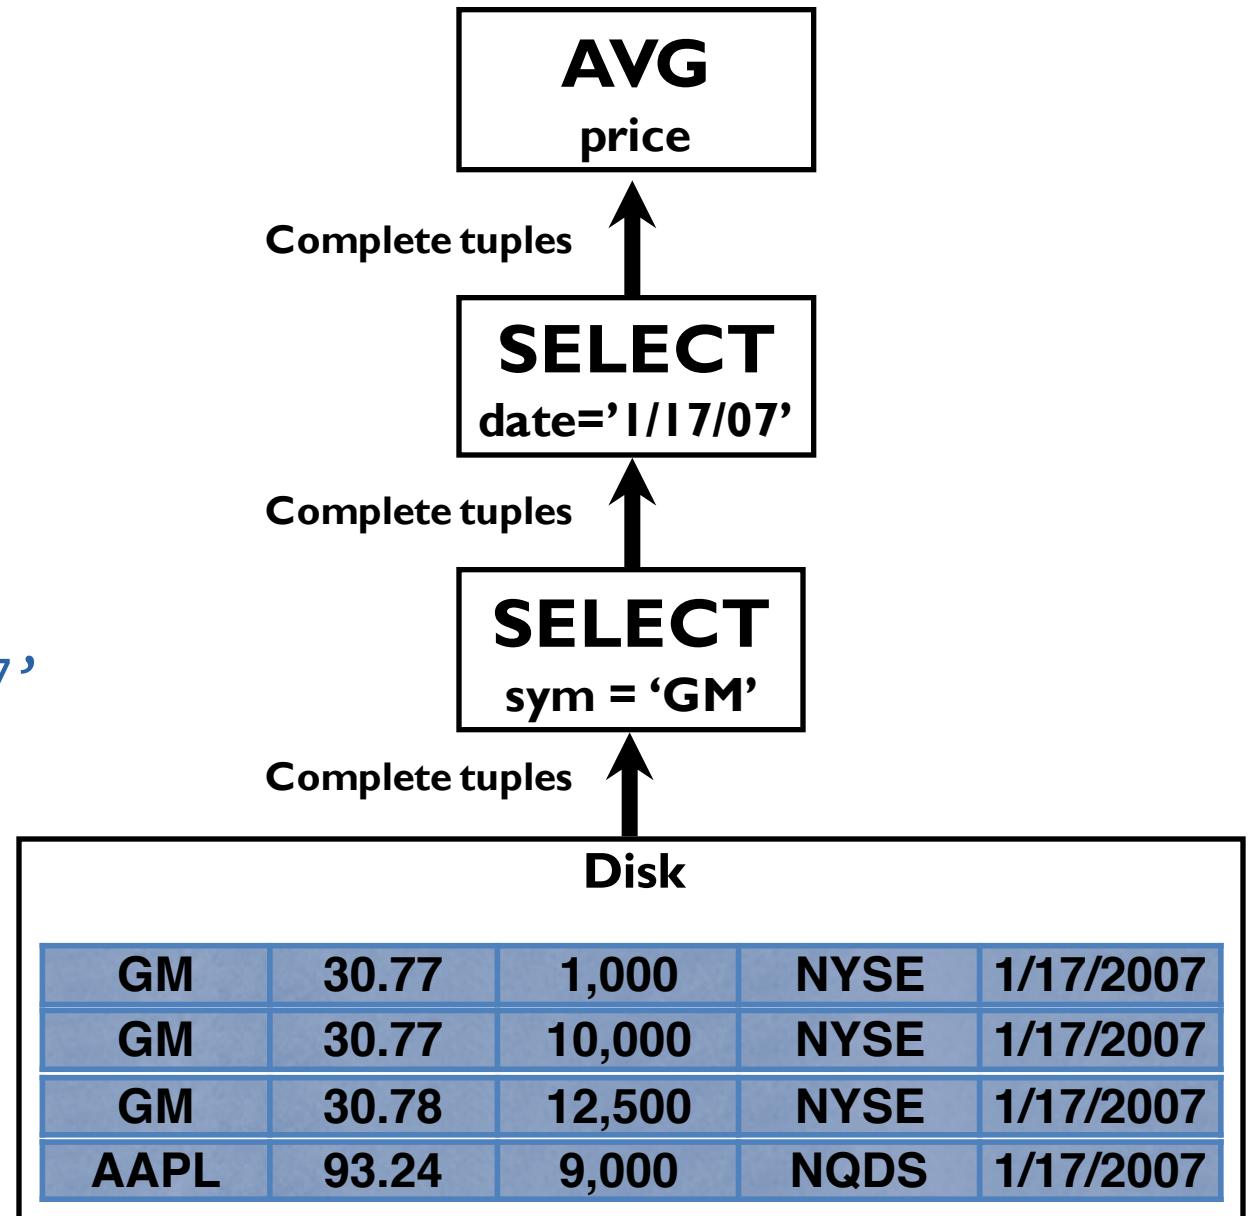
Compression better on single column

Execute on compressed data



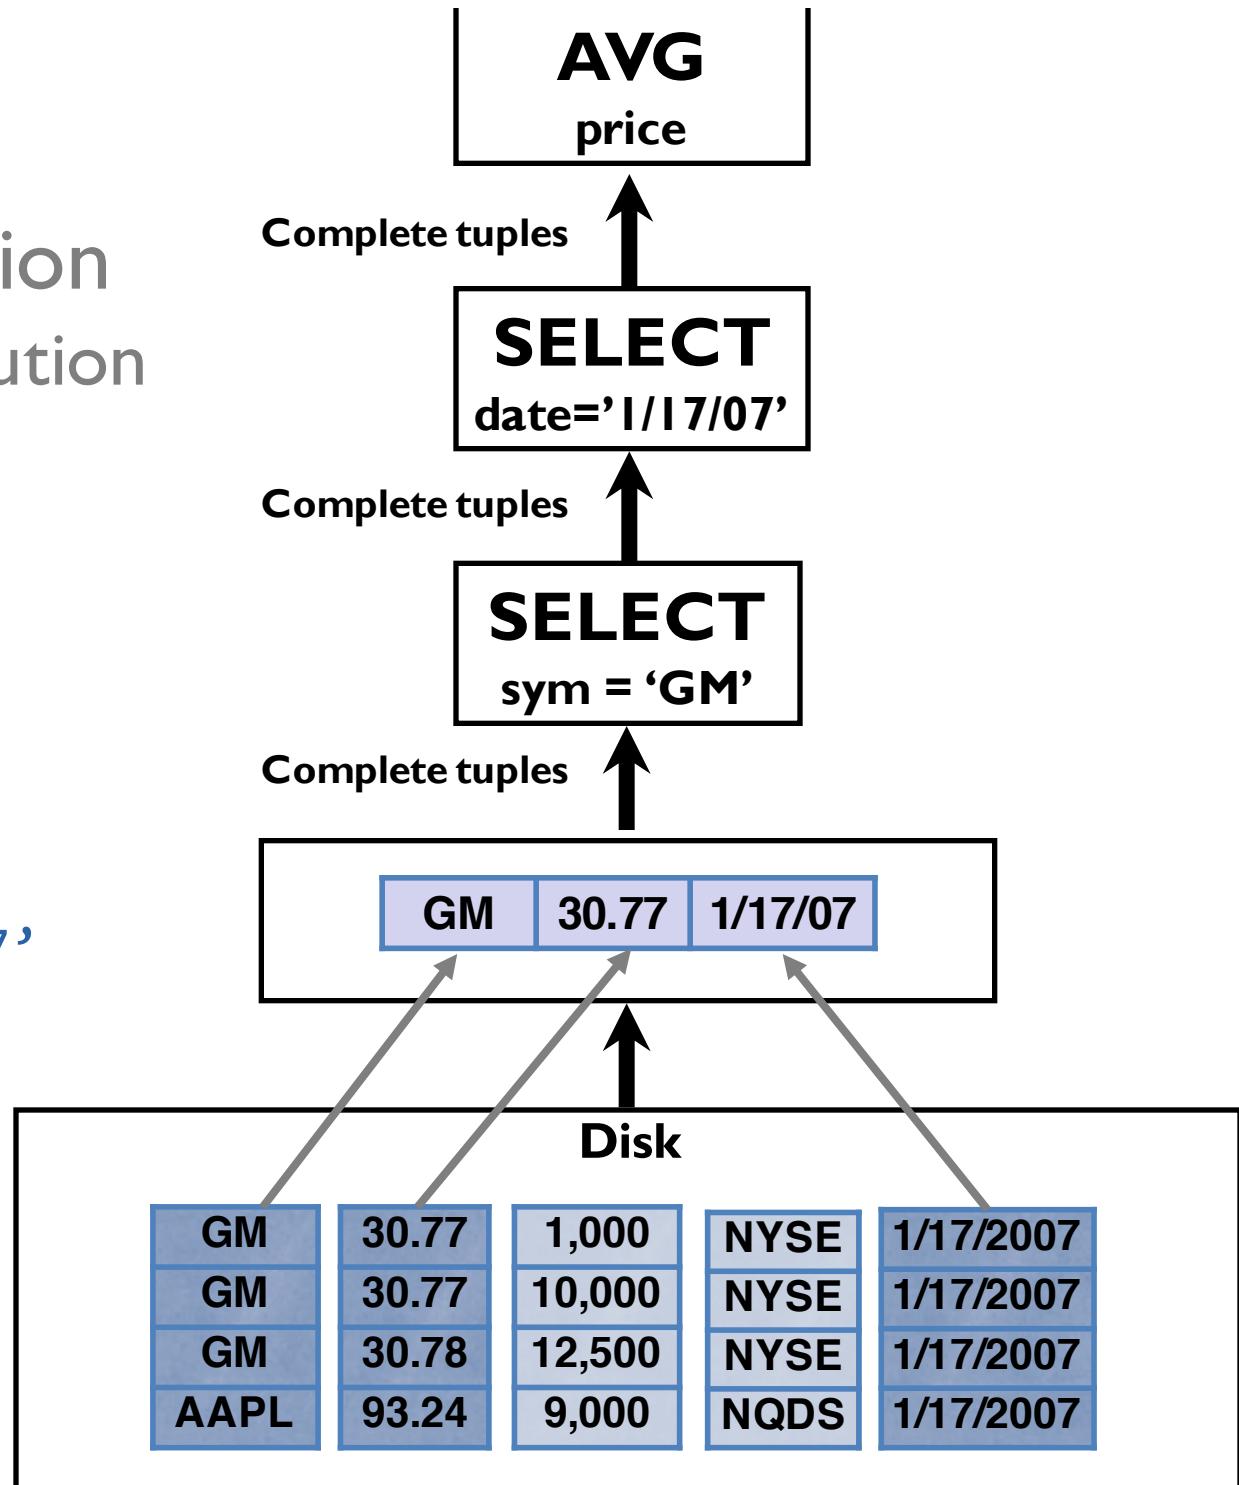
Traditional DBMS

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



Naïve: Early Materialization Row oriented execution

```
SELECT avg(price)  
FROM tickstore  
WHERE symbol = 'GM'  
AND date = '1/17/2007'
```

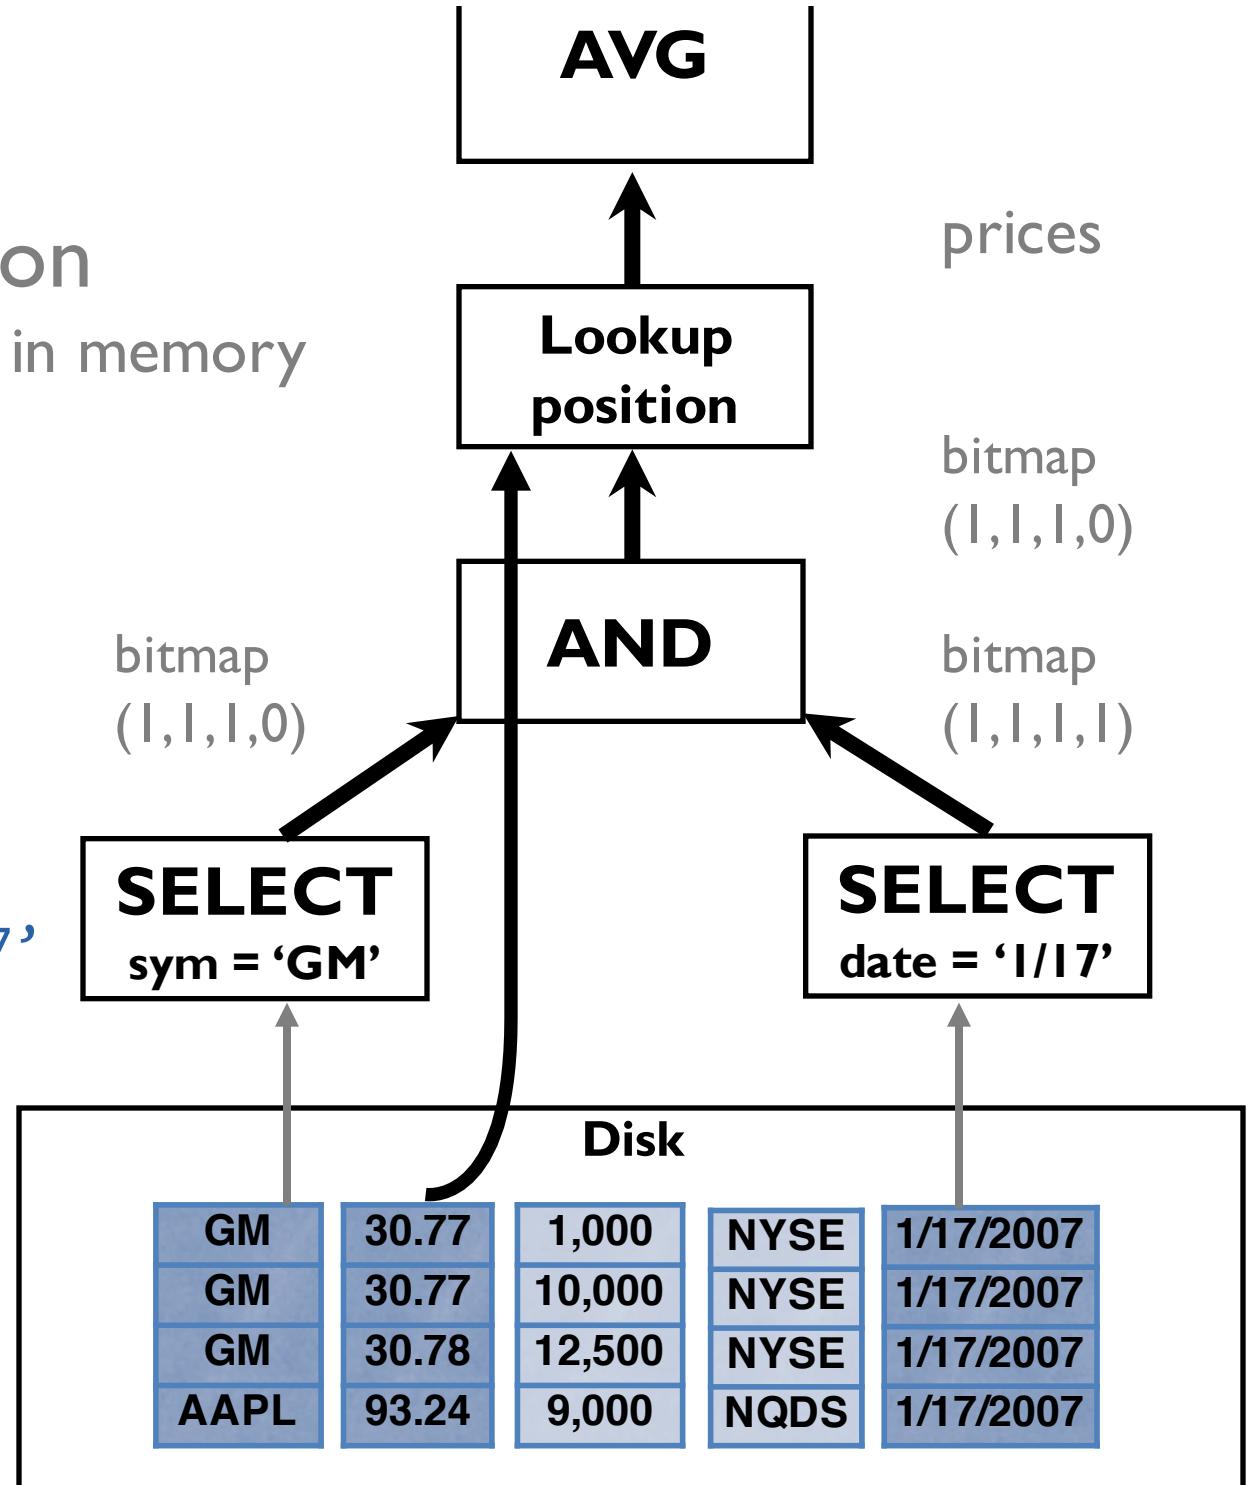


C-Store

Late Materialization

Much less data moving in memory

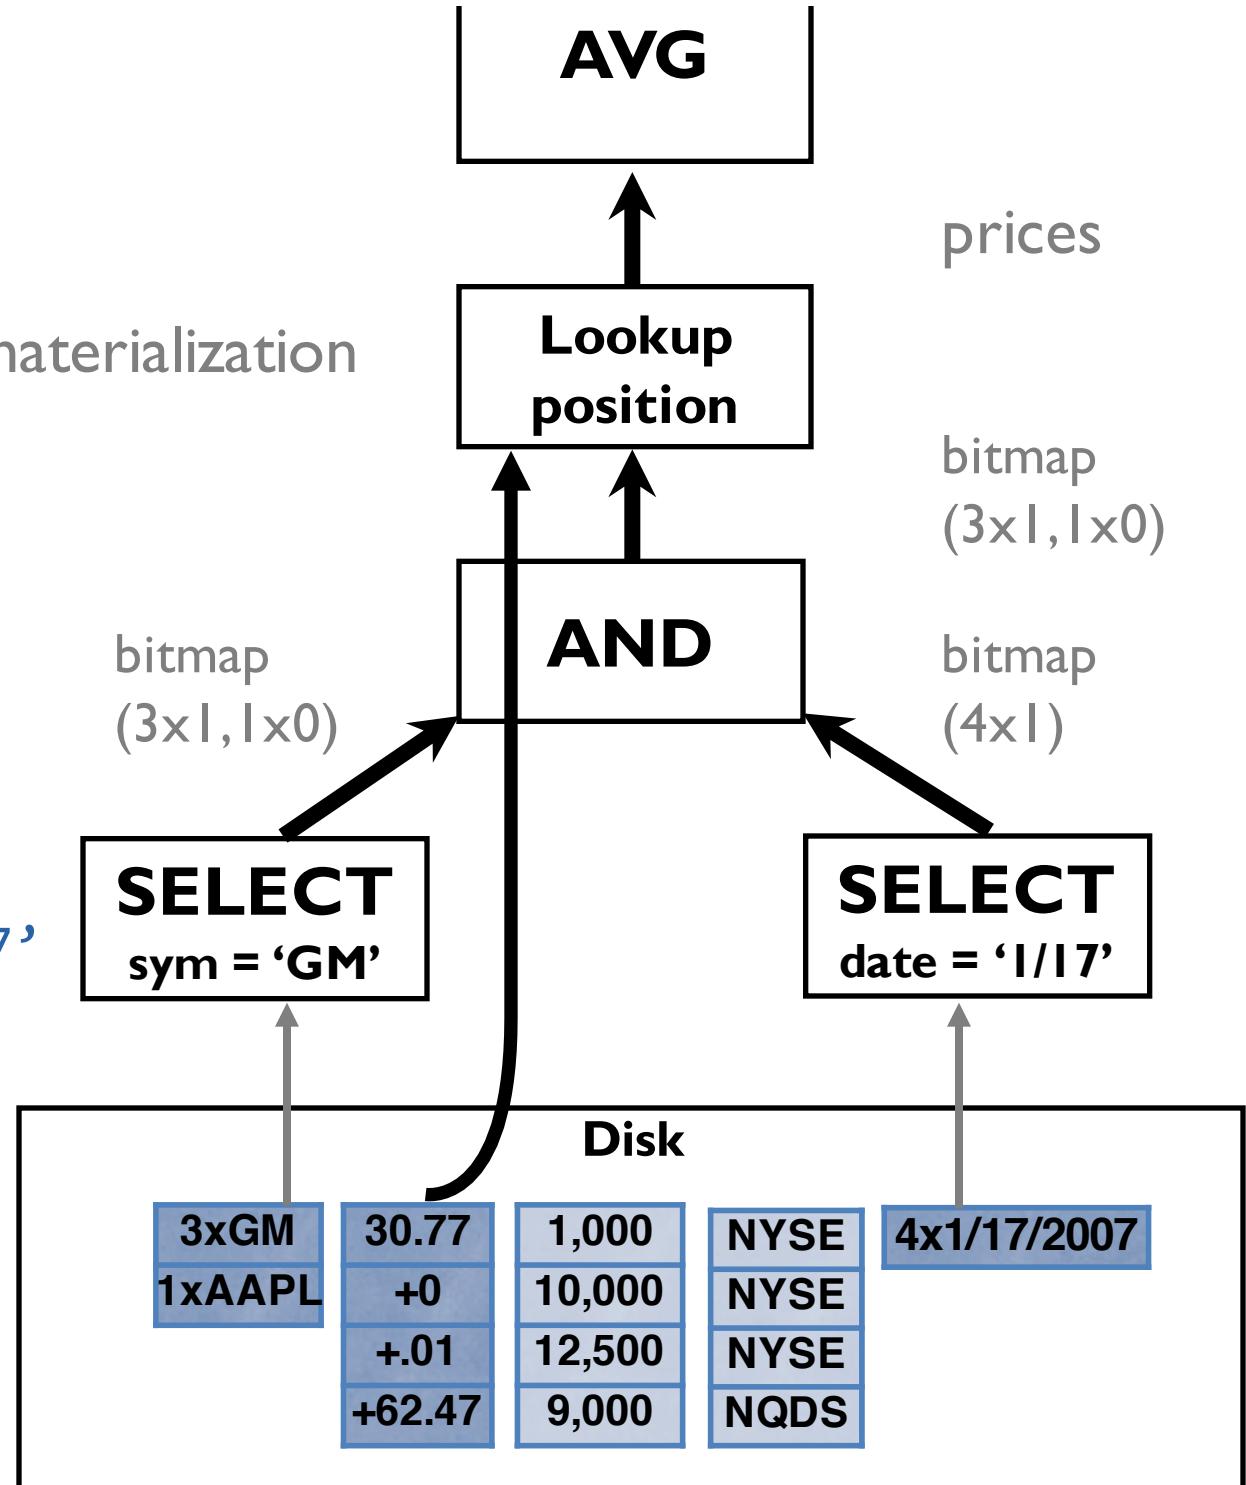
```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



C-Store Compression

Only possible w/ late materialization

```
SELECT avg(price)  
FROM tickstore  
WHERE symbol = 'GM'  
AND date = '1/17/2007'
```



Column Stores

Optimized for data warehouses

Store data by attribute/column rather than row

Compression

Compressed *query plan* execution

50-100x faster than row store

In-Memory DBMSes

Transaction-oriented apps

remove 1 unit from product

move 5 units from org 1 to org 2
(shopping carts, inventory)

Data stored in memory

Disk only used for check pointing

No concurrency

Active-active replication for fault-tolerance

Traditional Database

Indexes queries go faster

Concurrency queries go faster

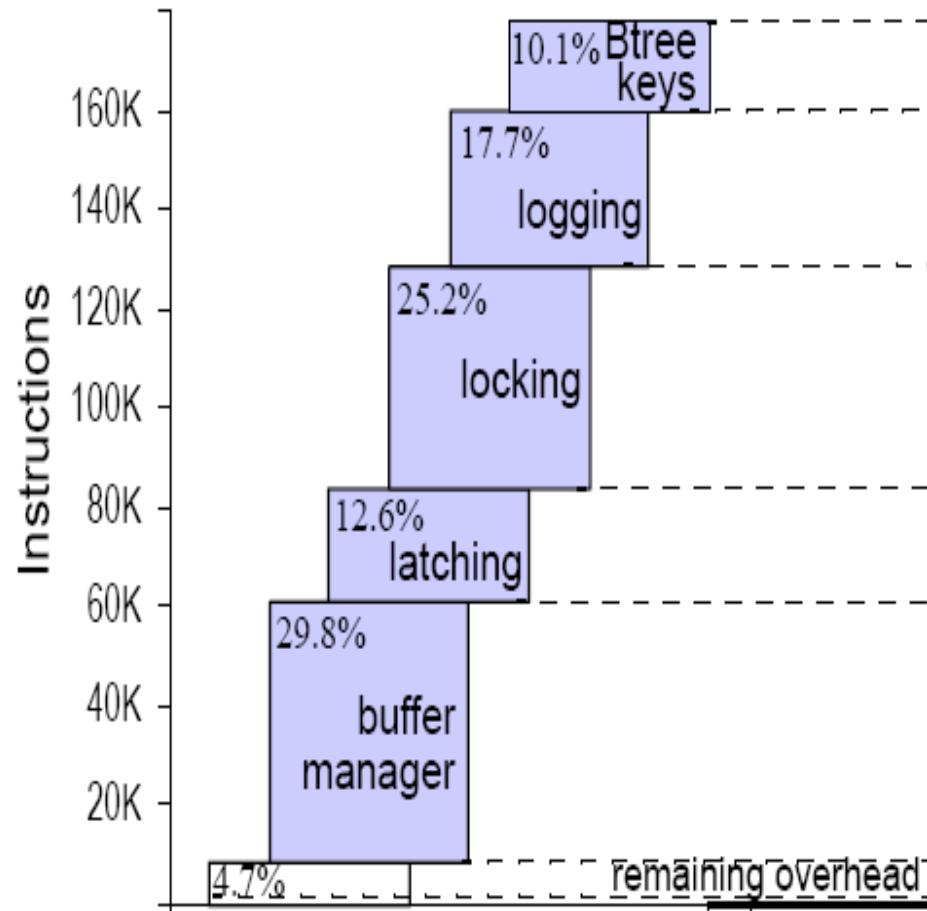
Locking serializability

Logging recovery

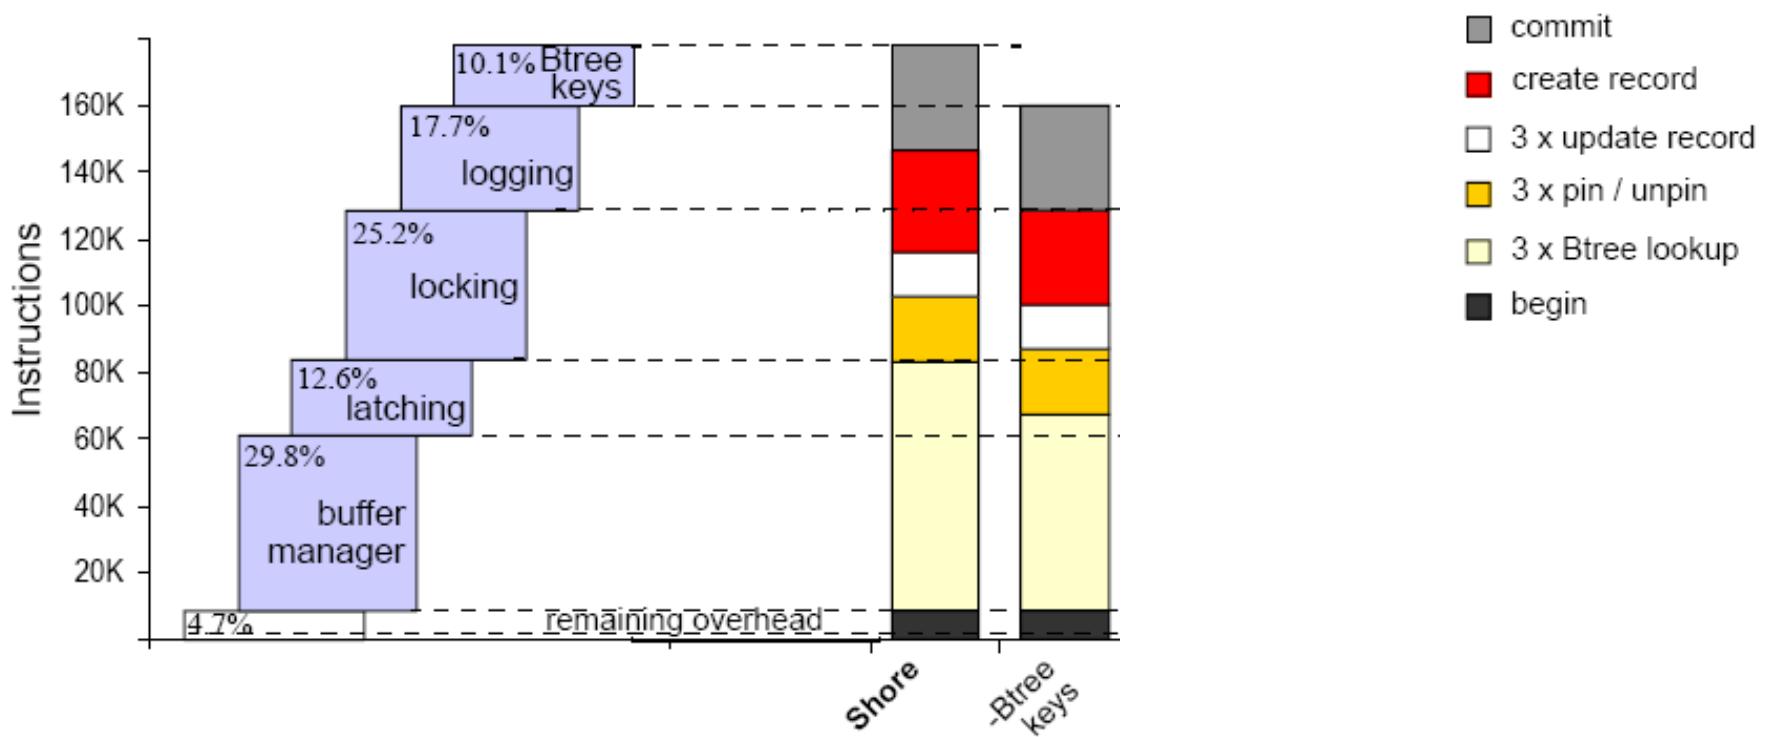
Buffer Manager manage pages in memory

Results after removing the components (in # instruction)

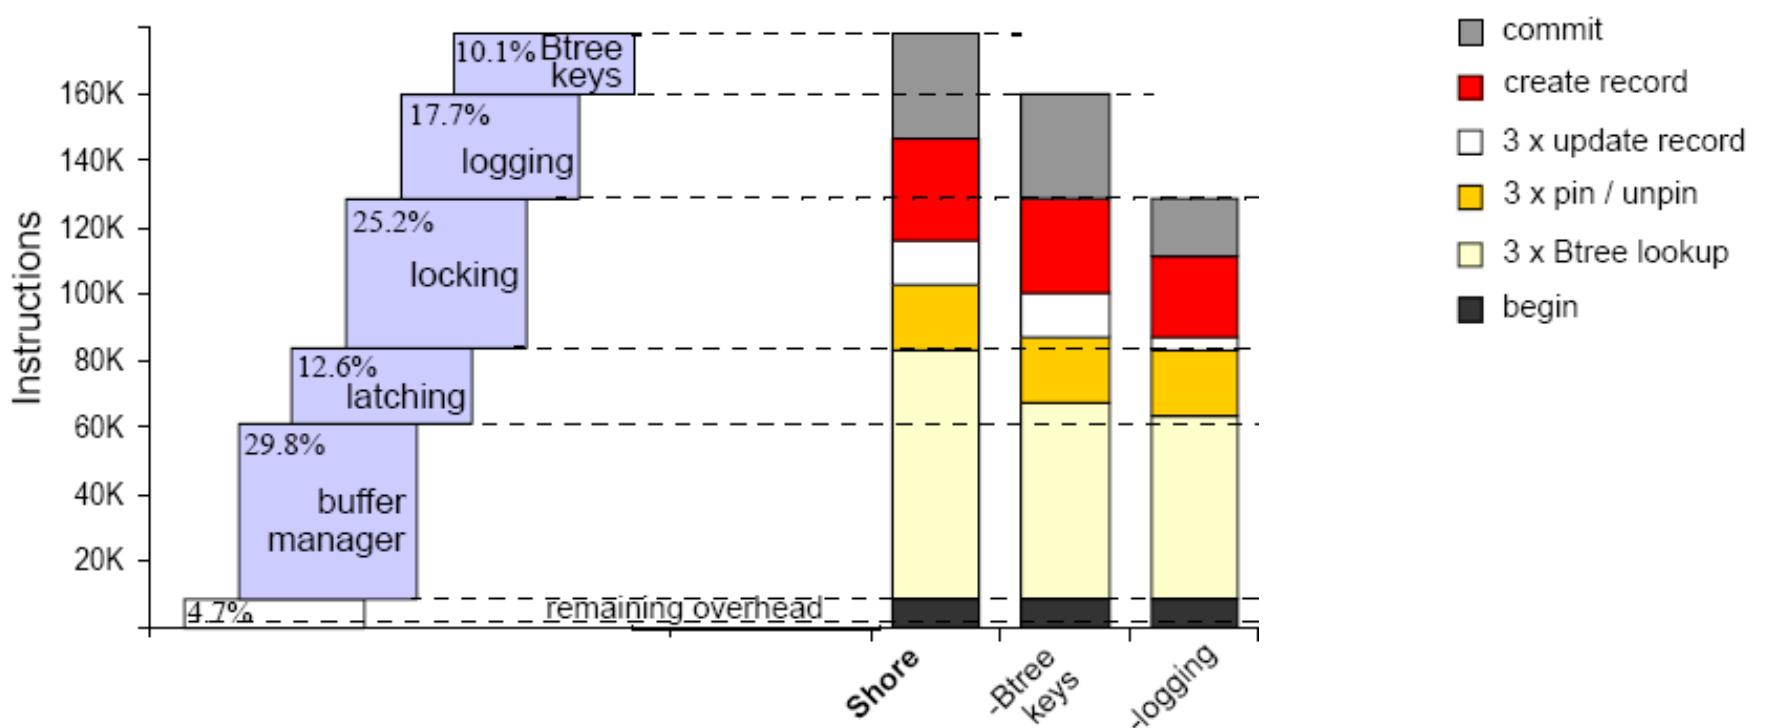
Instruction of useful work is only <2% of a memory resident DB



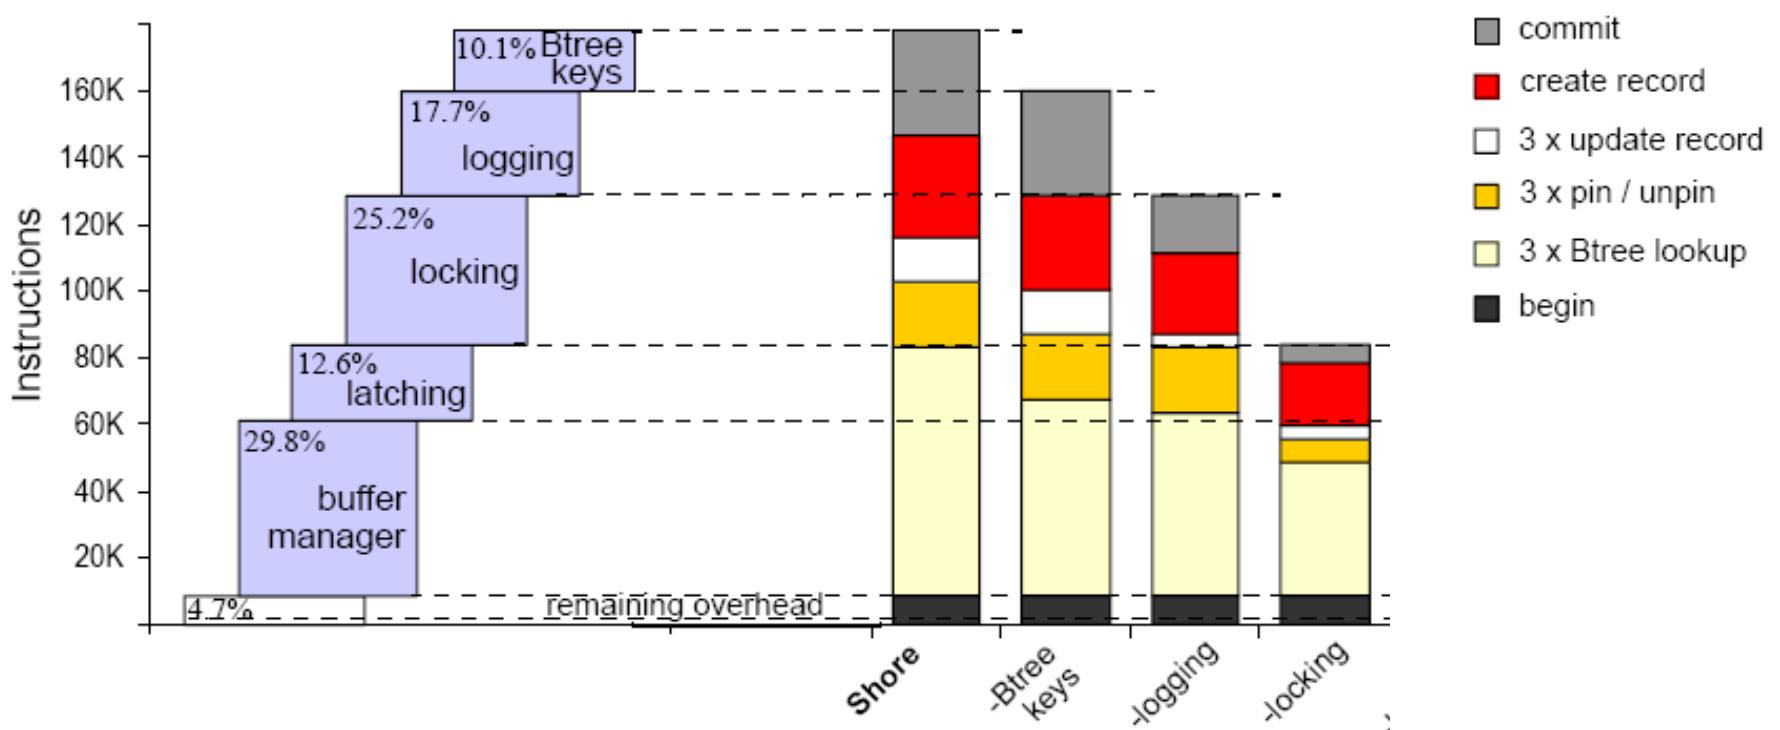
Effect of removing different components for payment (1/6)



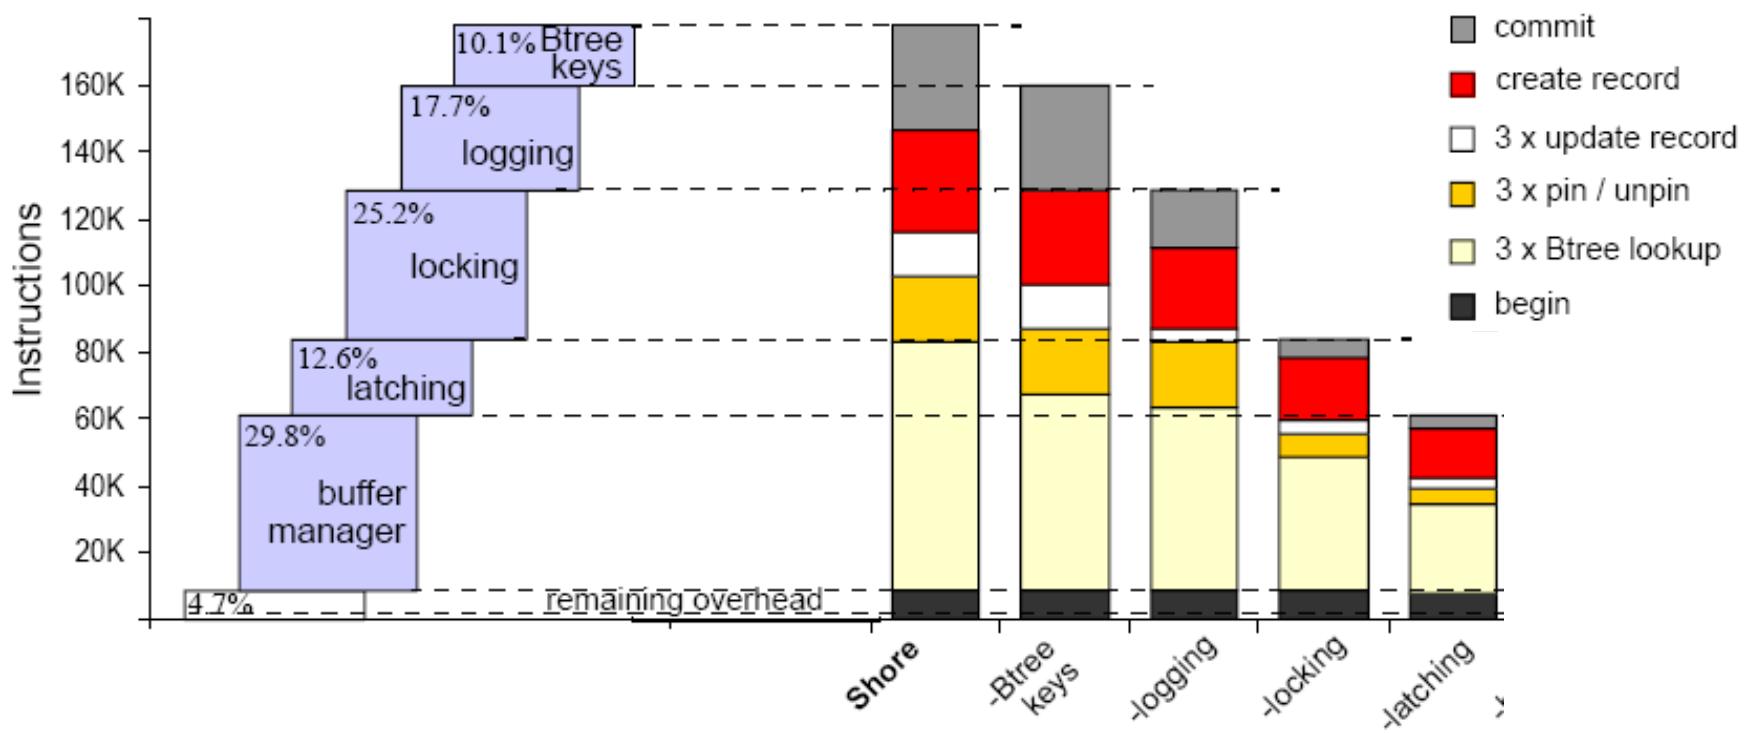
Effect of removing different components for payment (2/6)



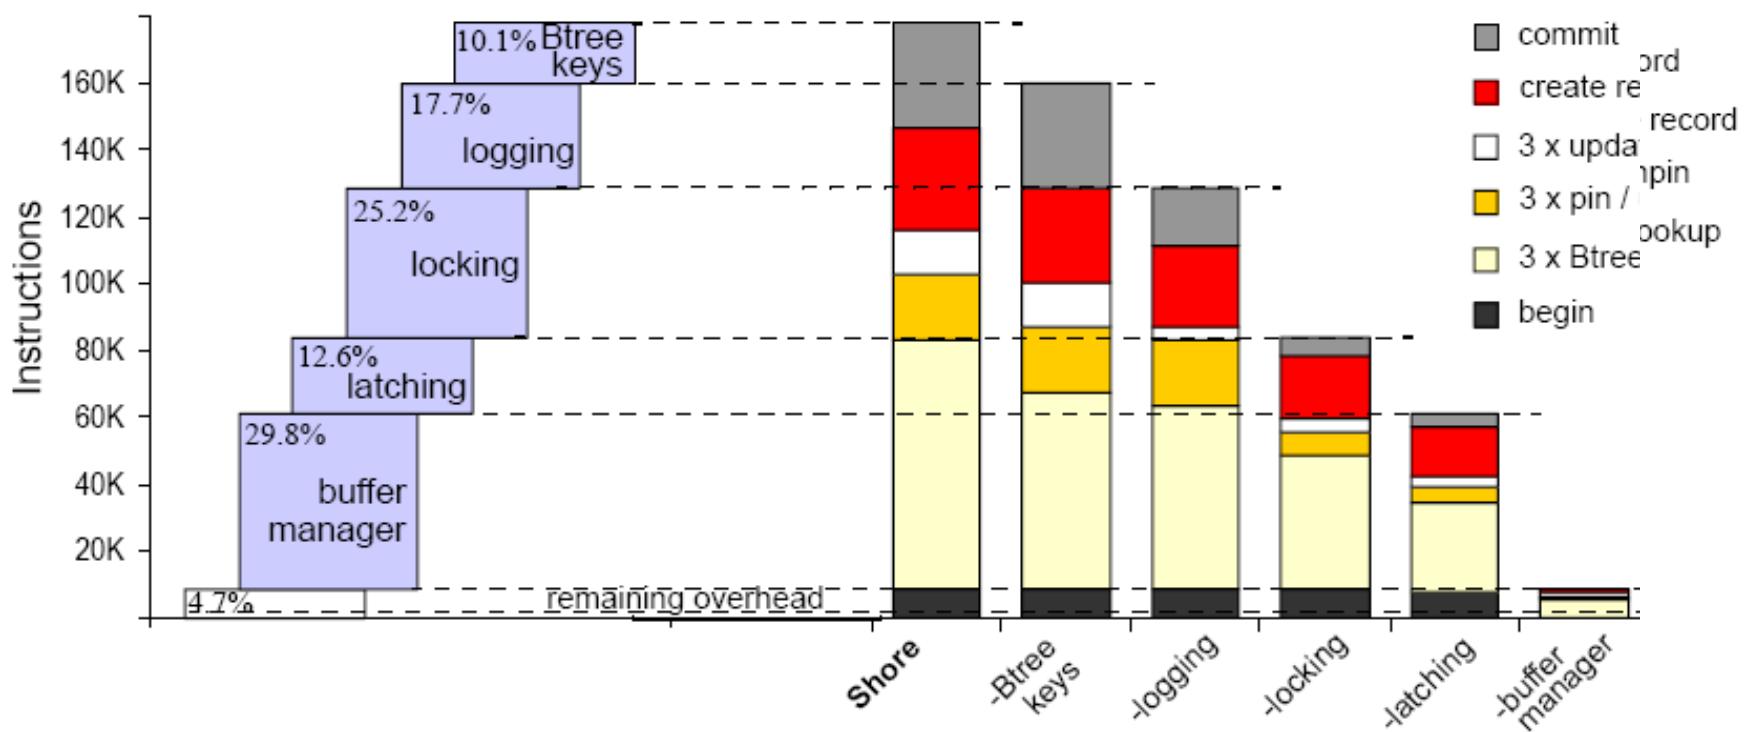
Effect of removing different components for payment (3/6)



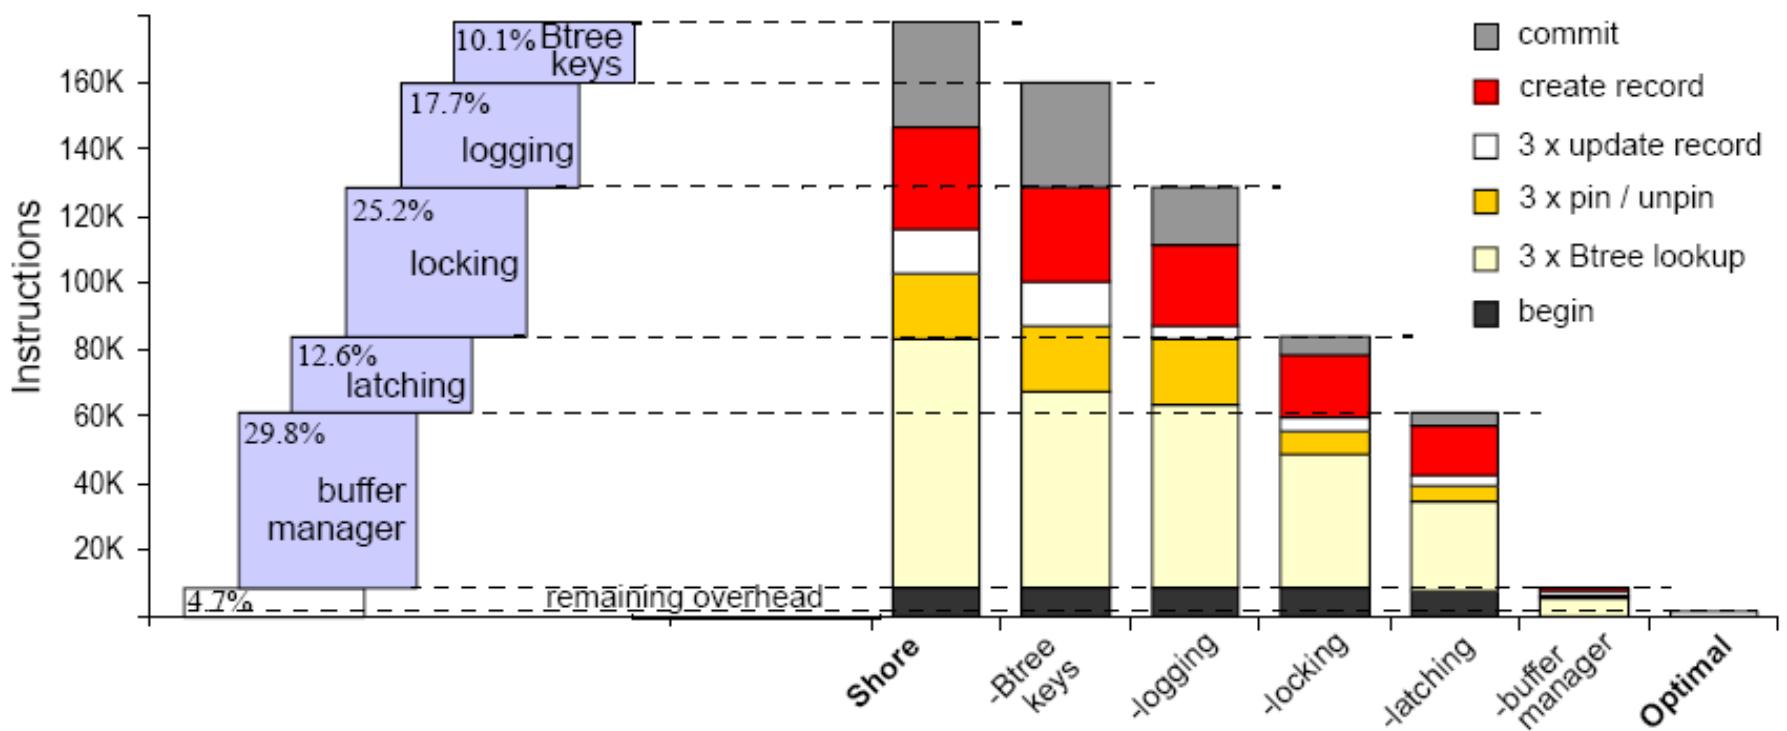
Effect of removing different components for payment (4/6)



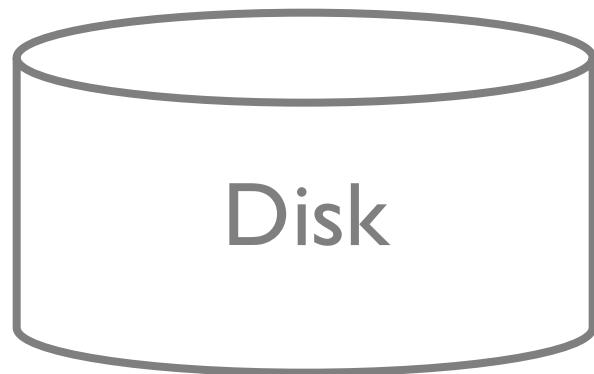
Effect of removing different components for payment (5/6)



Effect of removing different components for payment (6/6)



Barebones
Executor

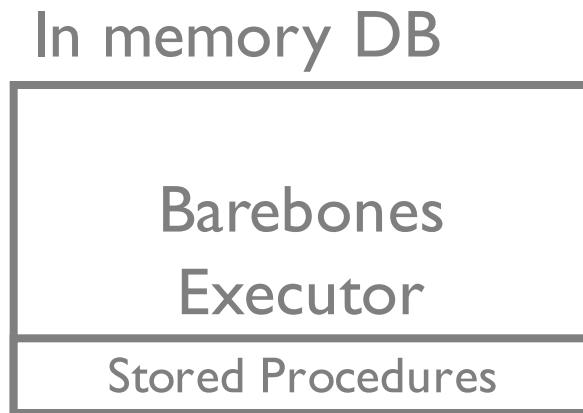


In memory DB

Barebones
Executor

What about
Parsing
Concurrency?
Recovery?

**What about
Parsing
Concurrency?
Recovery?**



Procedure:
`p1 = SELECT cost
 FROM fact_table
 WHERE id = ?`

Query:
`p1(10)`

In memory DB

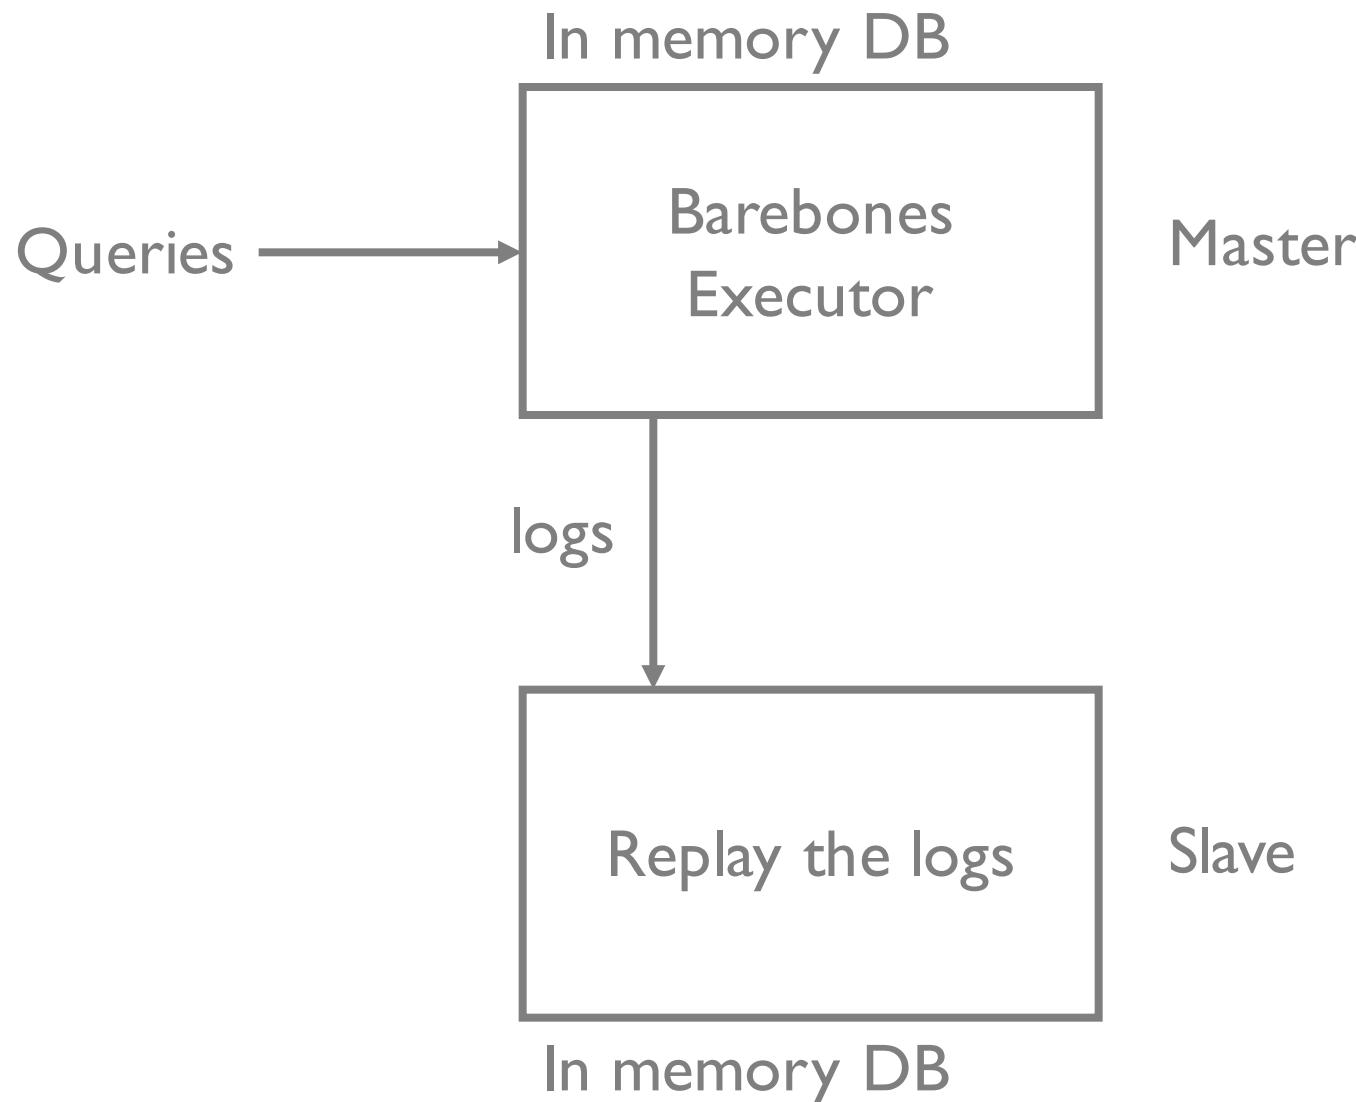
Barebones
Executor

What about
Parsing
~~Concurrency?~~

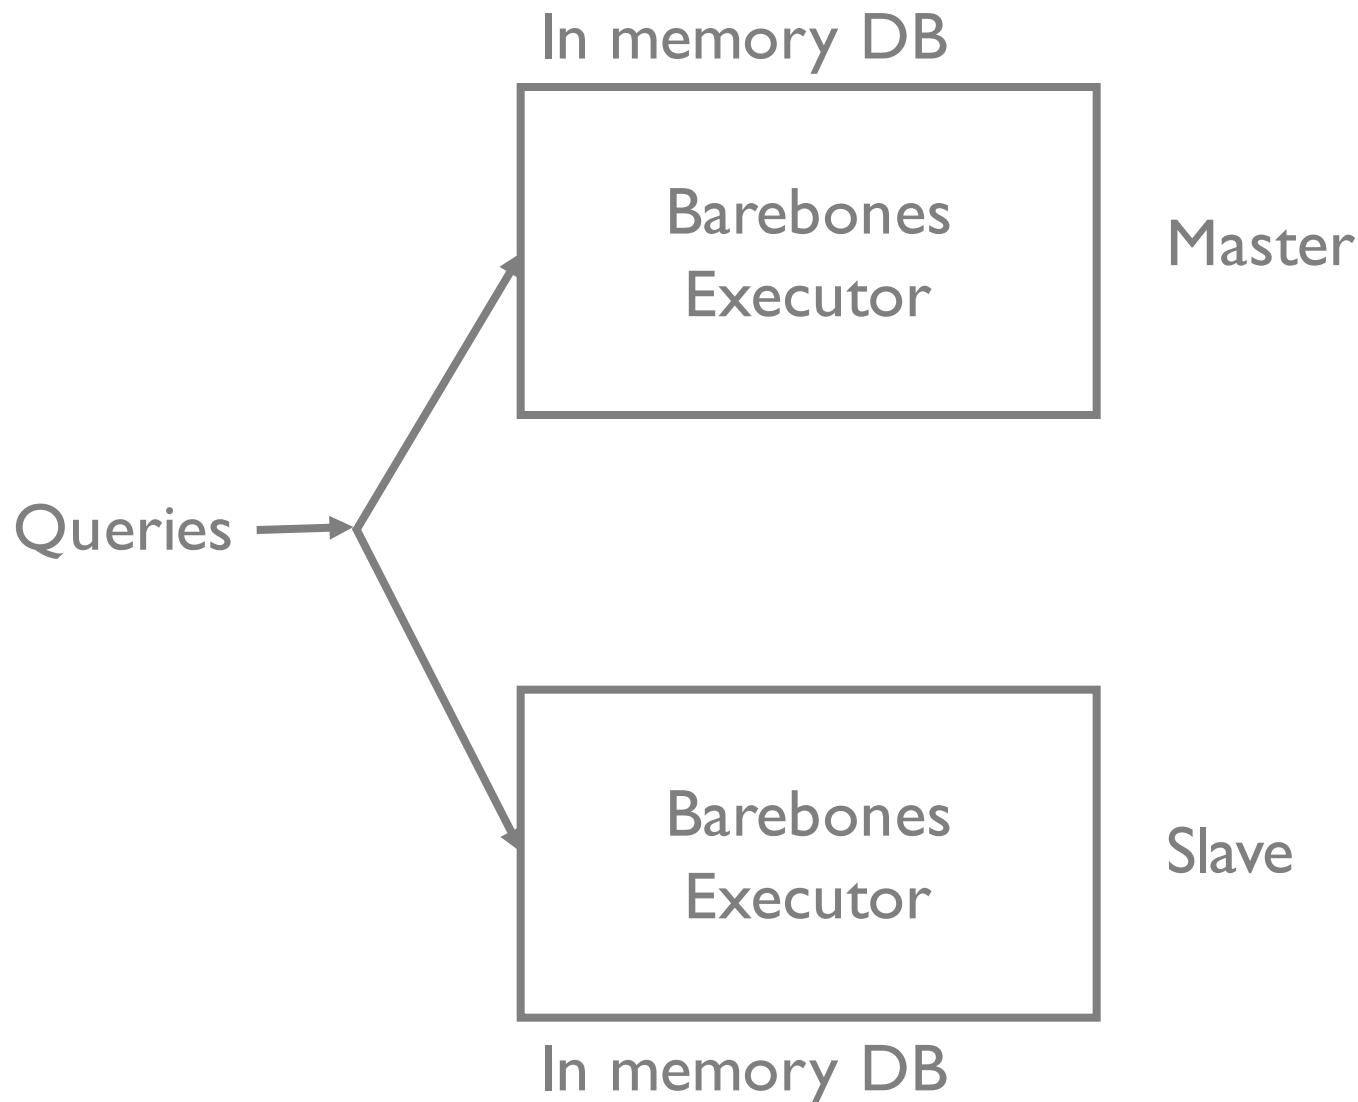
no buffer manager, no concurrency, no locks

Recovery?

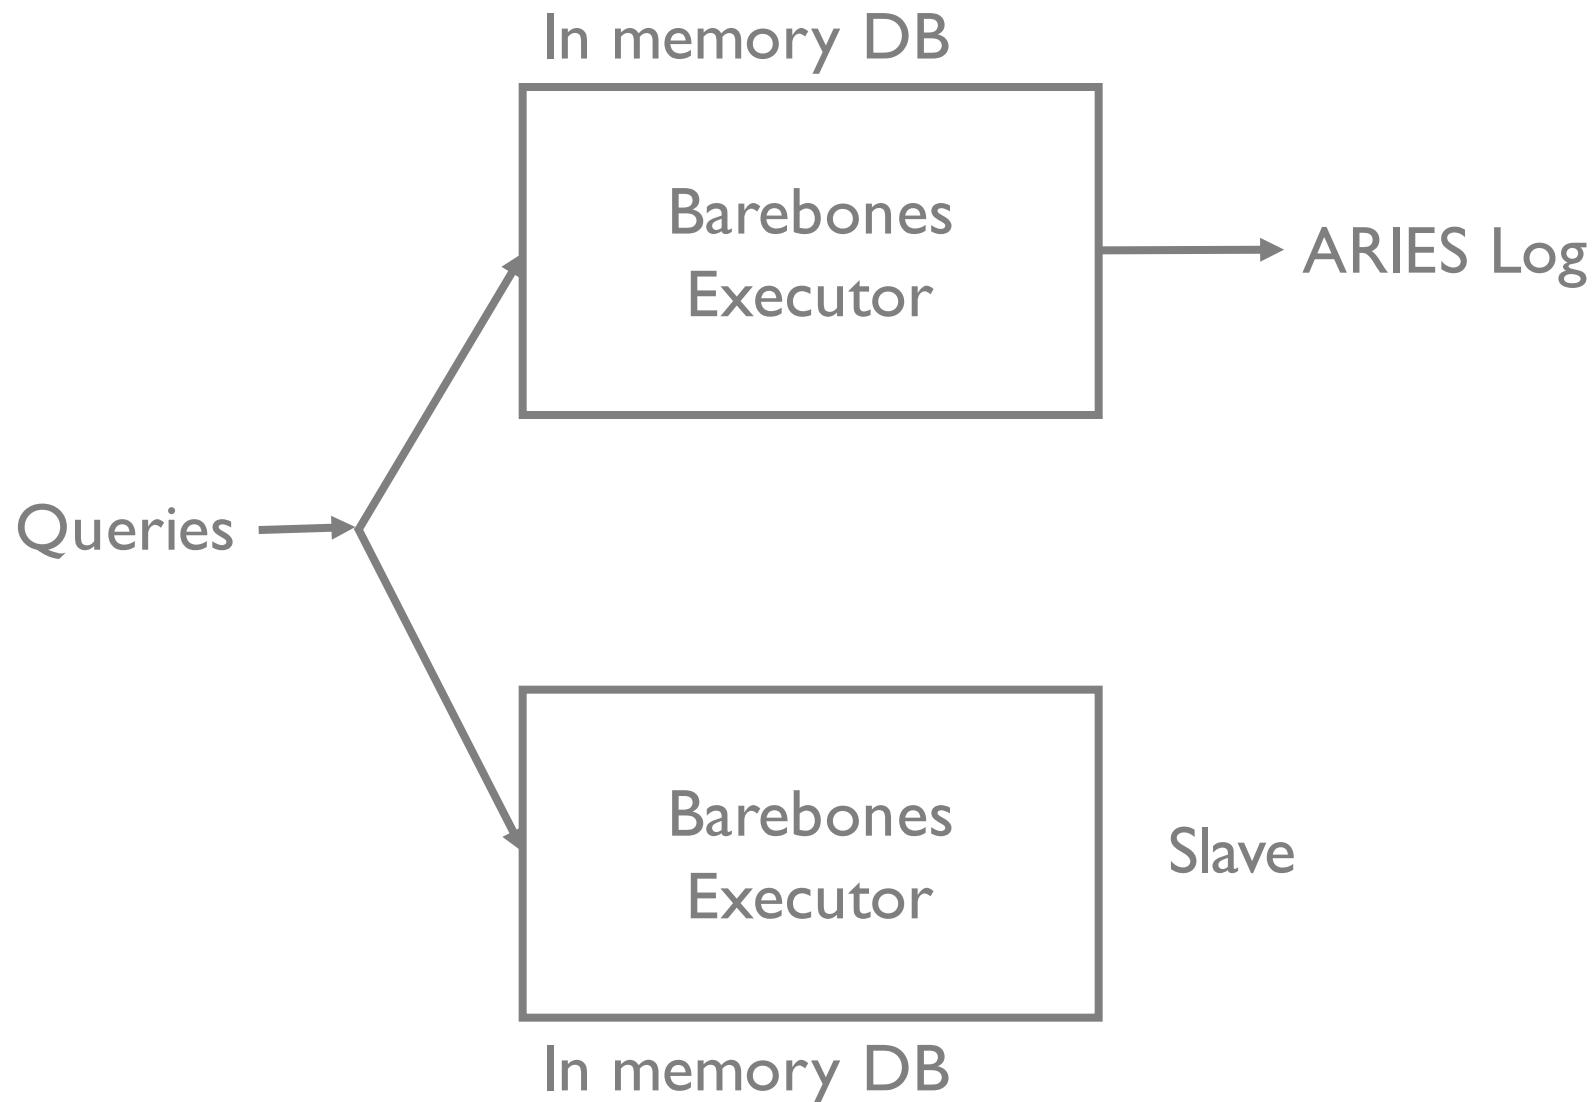
Log Shipping



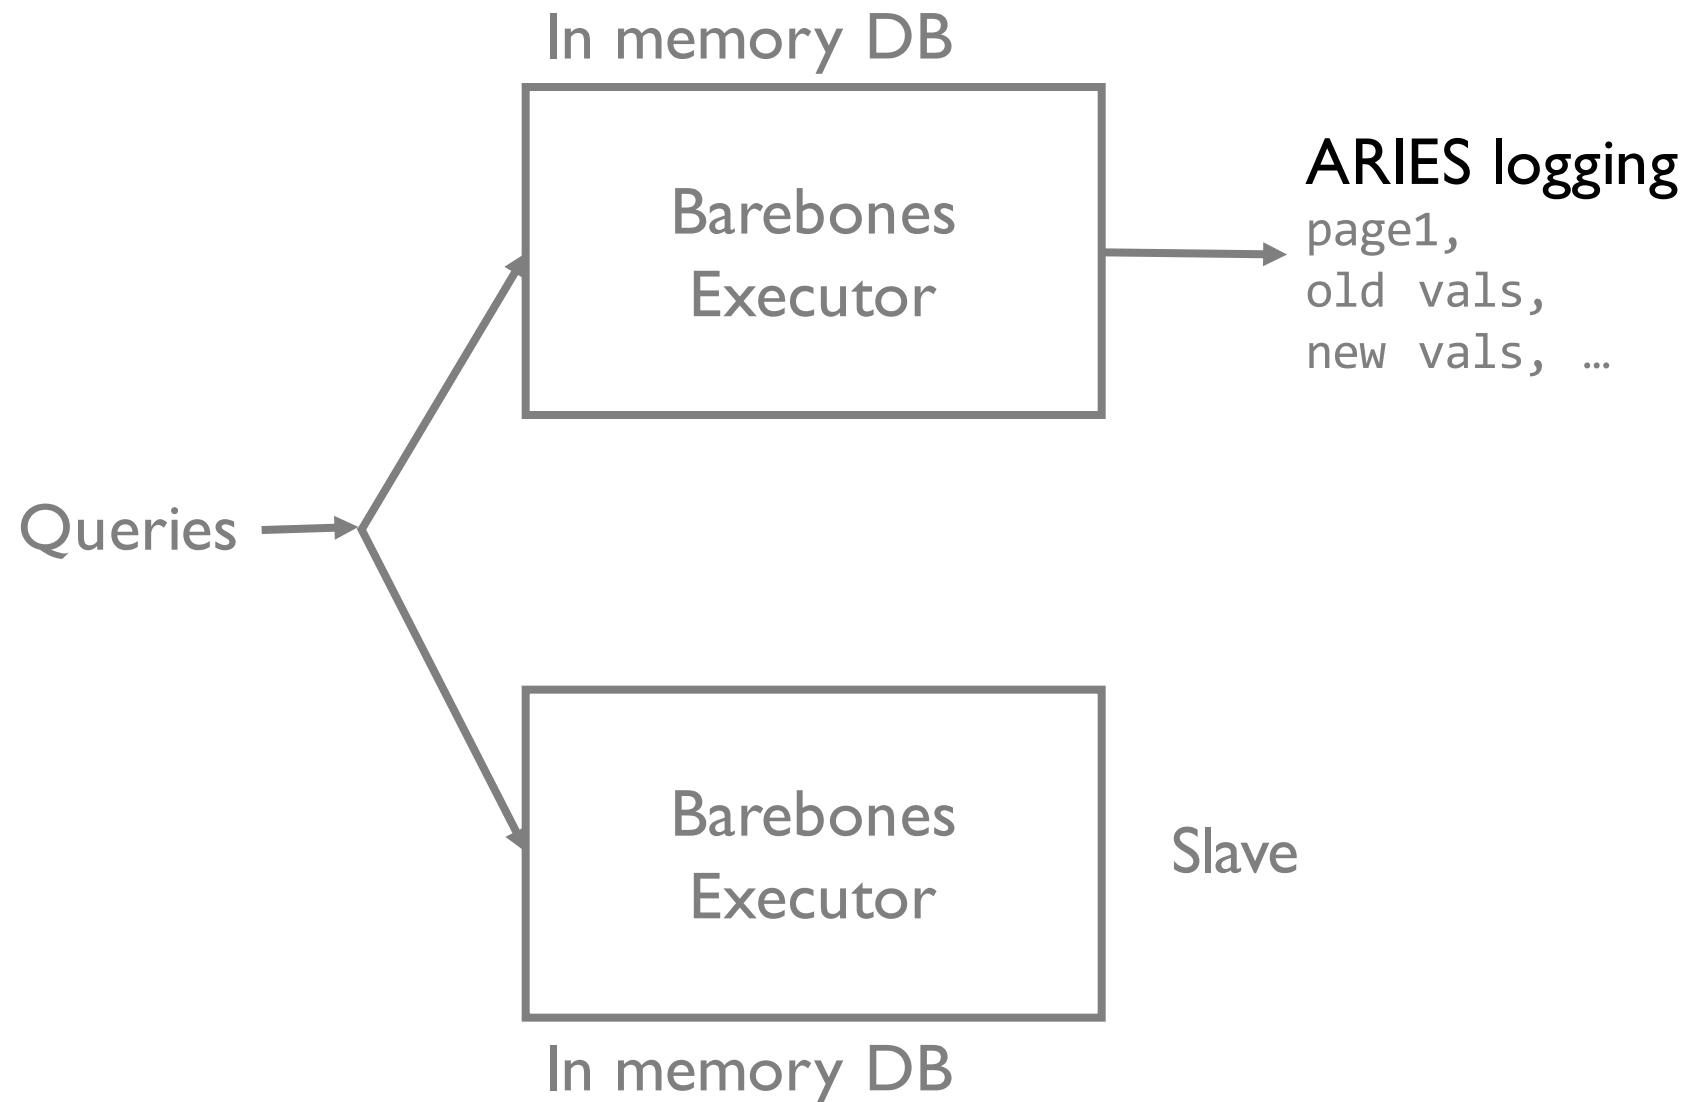
Active-Active



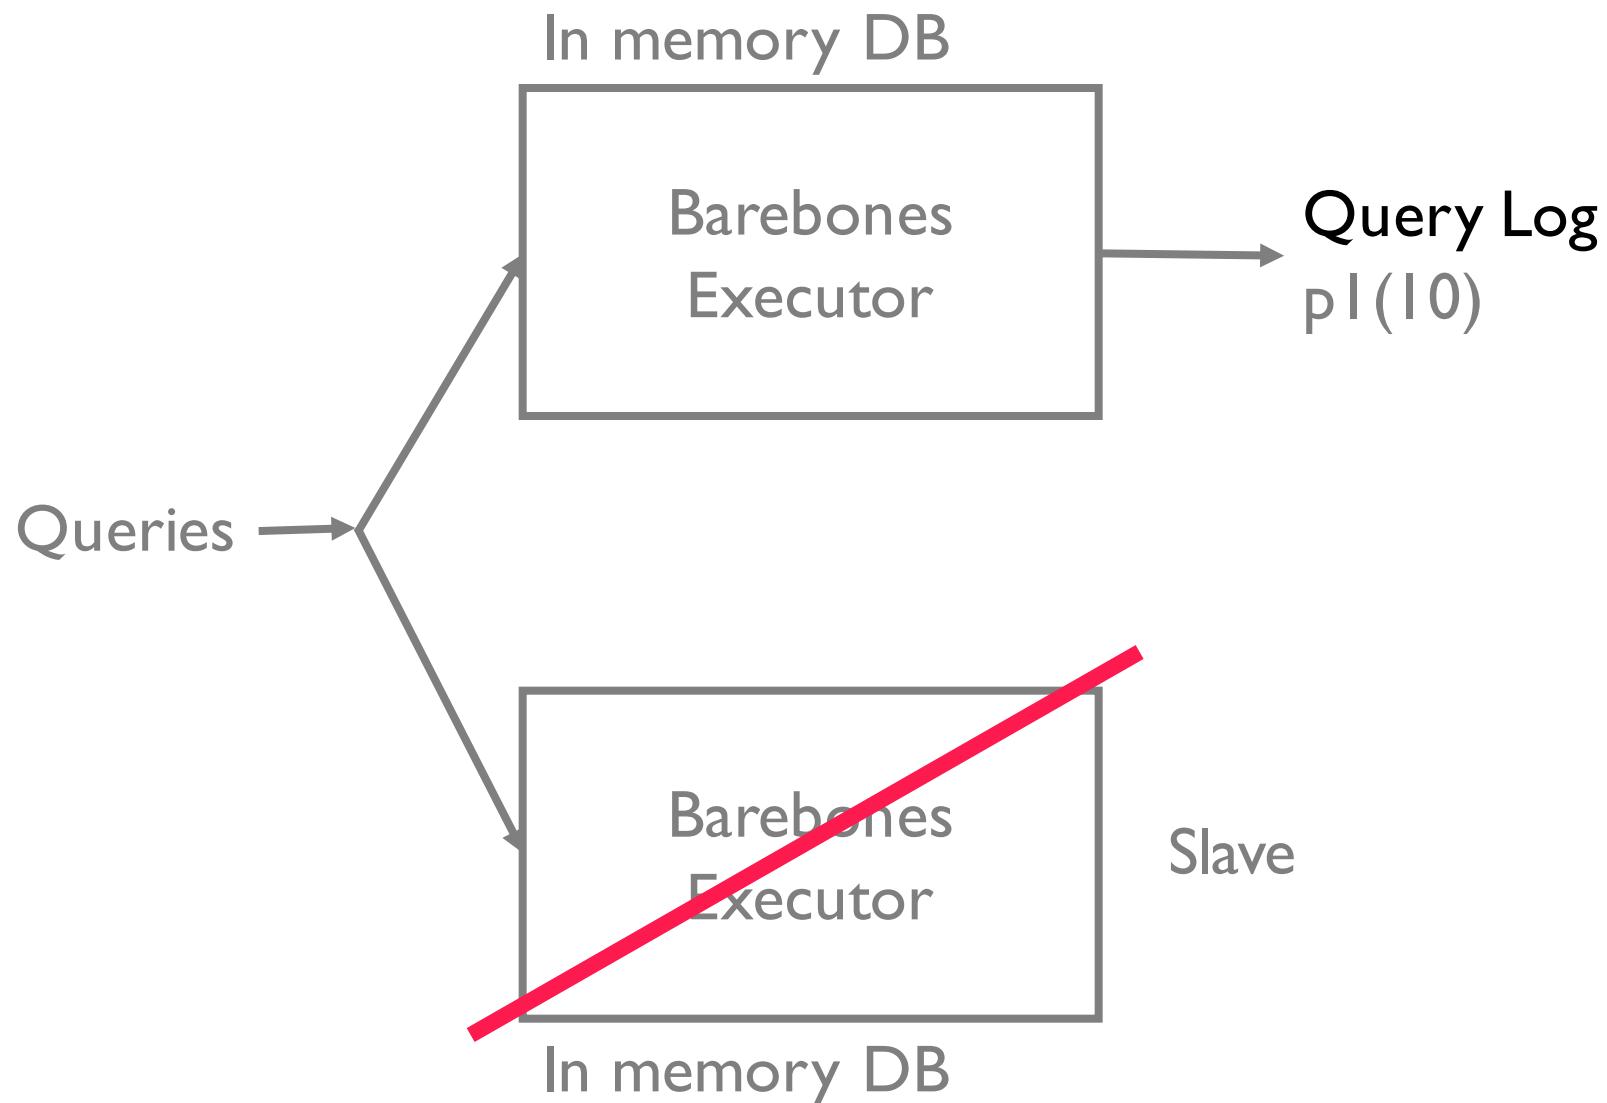
Recovery in Active-Active



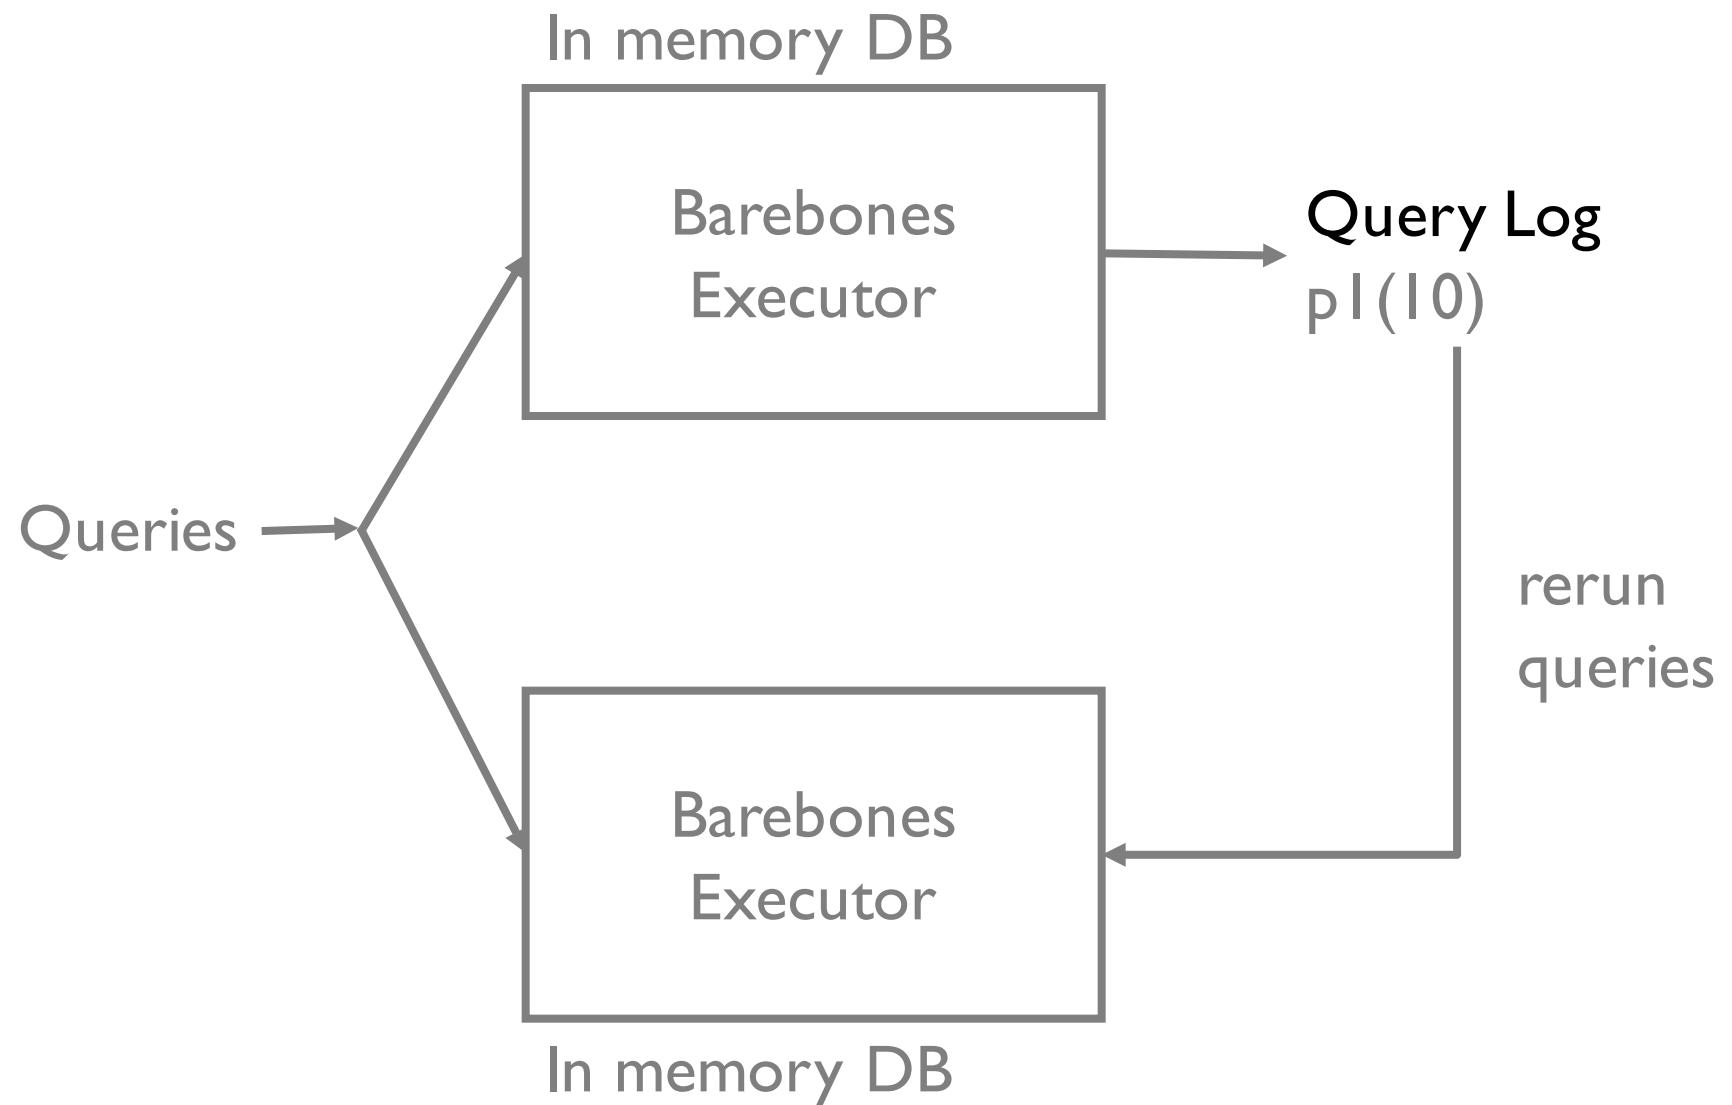
Recovery in Active-Active



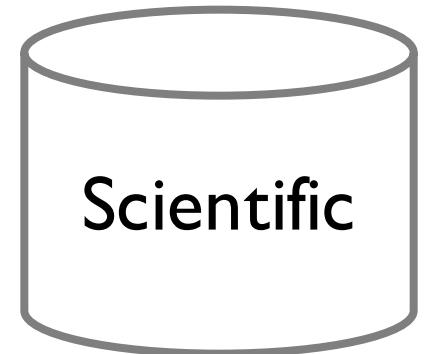
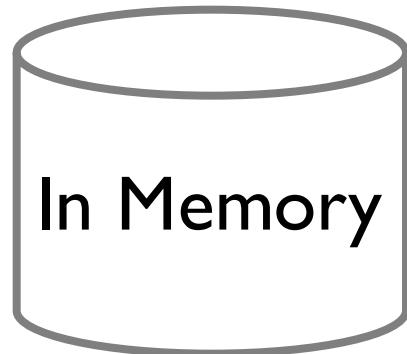
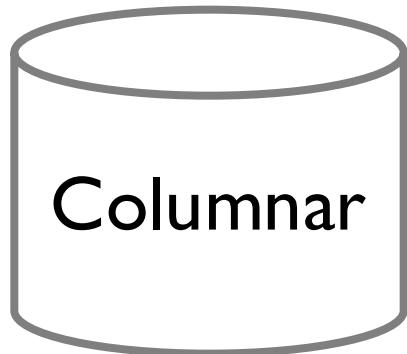
Recovery in Active-Active



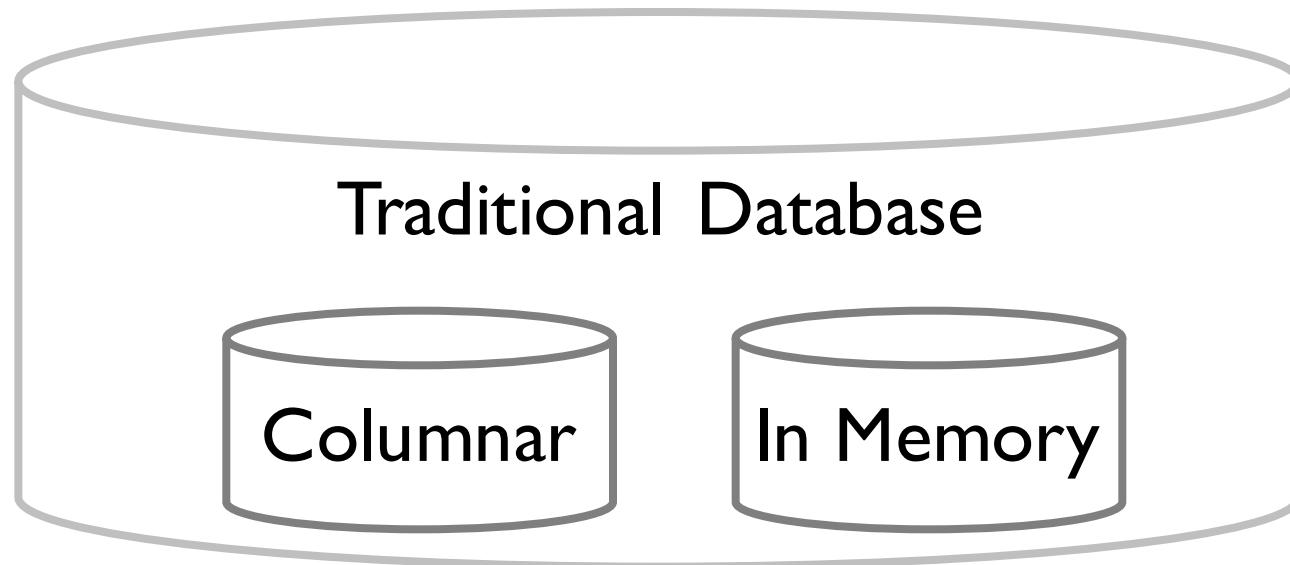
Recovery in Active-Active



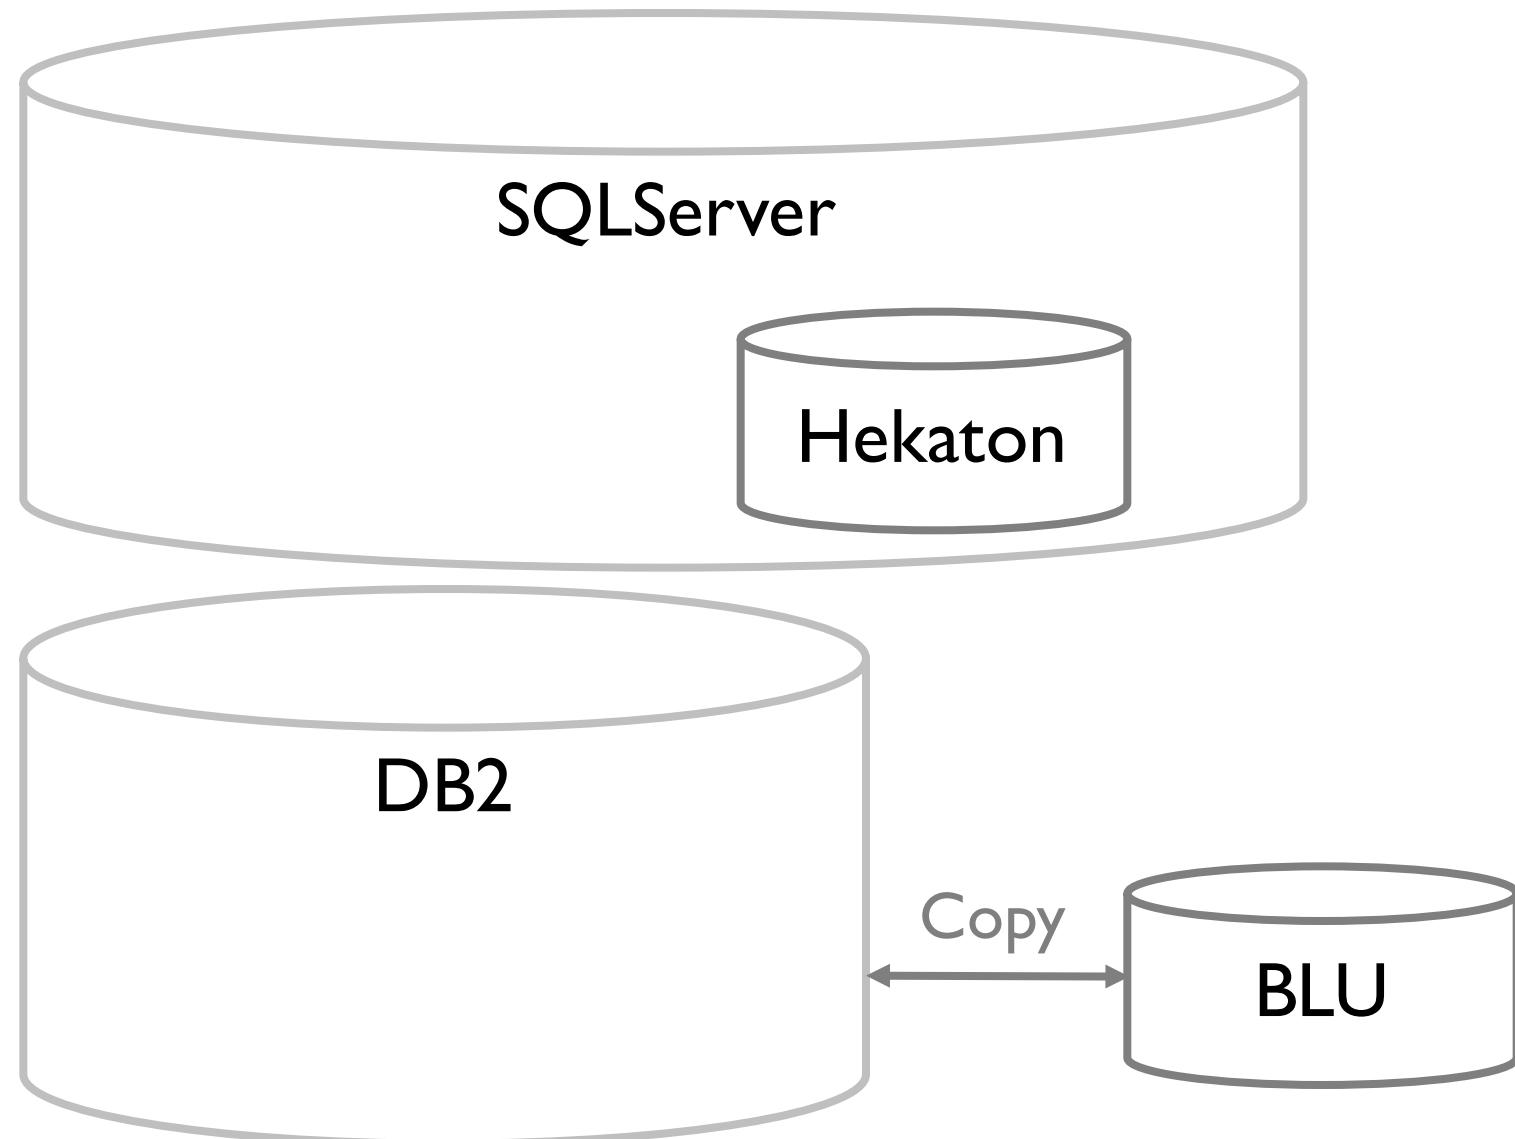
One Size Does Not Fits All



One Size Does Not Fit All



One Size Does Not Fits All



OK Let's Step Back

Discussed the *how*:

how to model data needed by an application

how to query databases

how databases execute queries

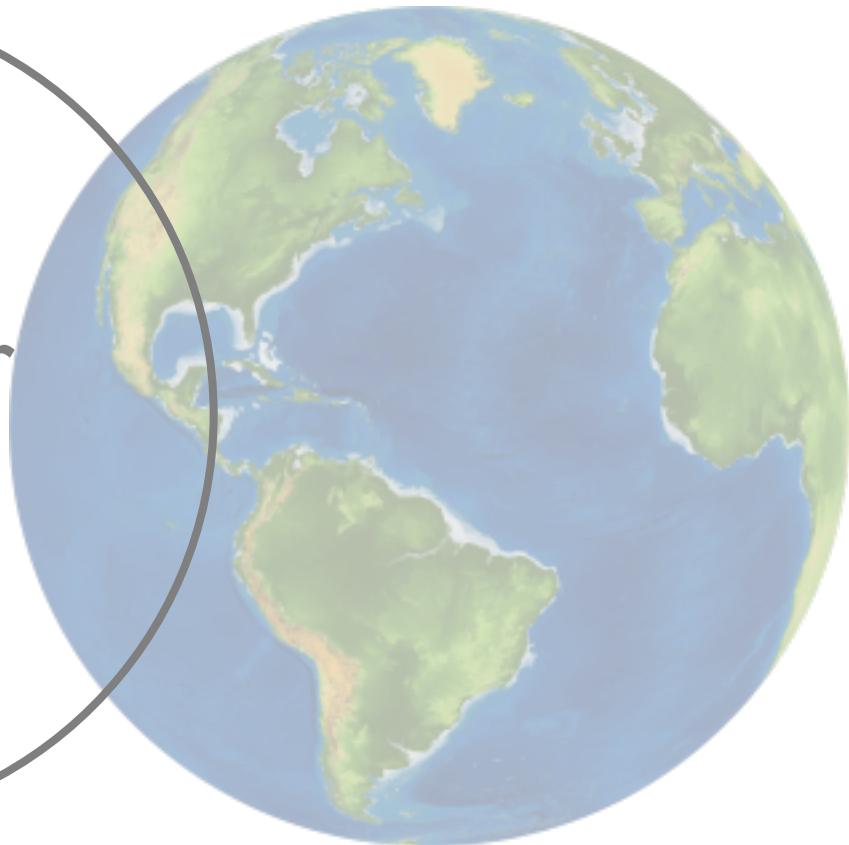
how databases run fast, correctly

To what end?

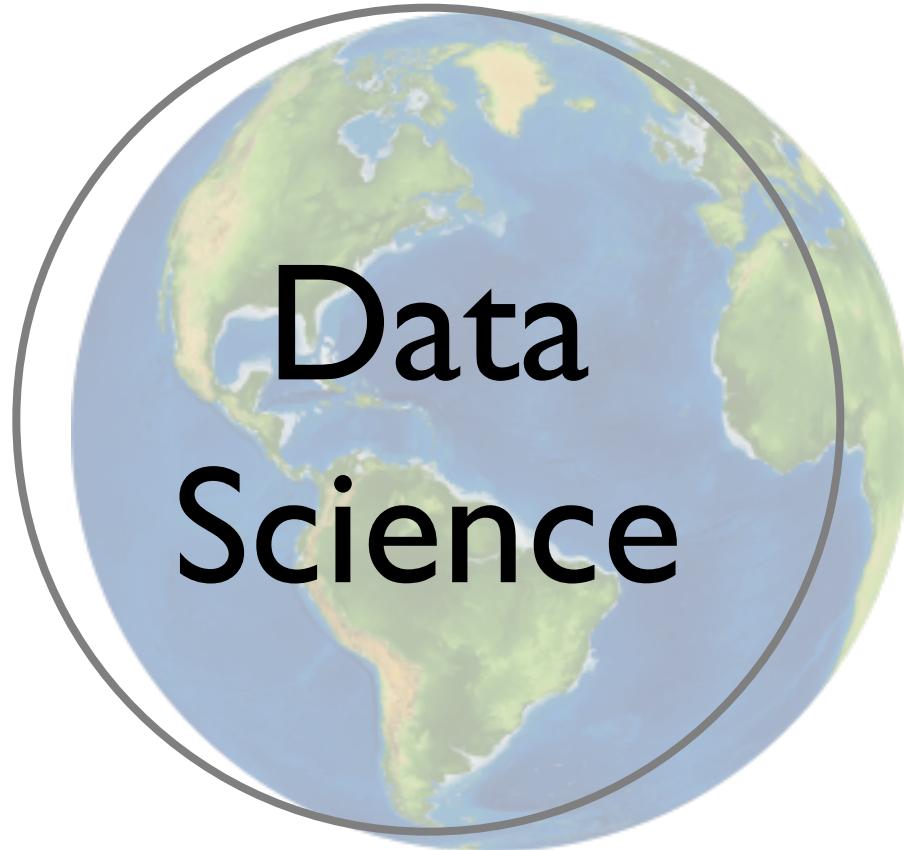
Computer
Science



Computer
Science







**Data is crucial to
our lives as individuals and as a society**

Thanks!

Please fill out courseworks survey

Role of 4111 in DSI curriculum